

TCP/IP in hardware using SME

Mark Jan Jacobi & Jan Meznik

KU

September 18, 2019

TCP/IP in hardware using SME

2019-09-18

TCP/IP in hardware using SME

Mark Jan Jacobi & Jan Meznik

KU

September 18, 2019

Mark siger introduktion og 2-3 sætninger "abstrakt"

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18

TCP/IP in hardware using SME

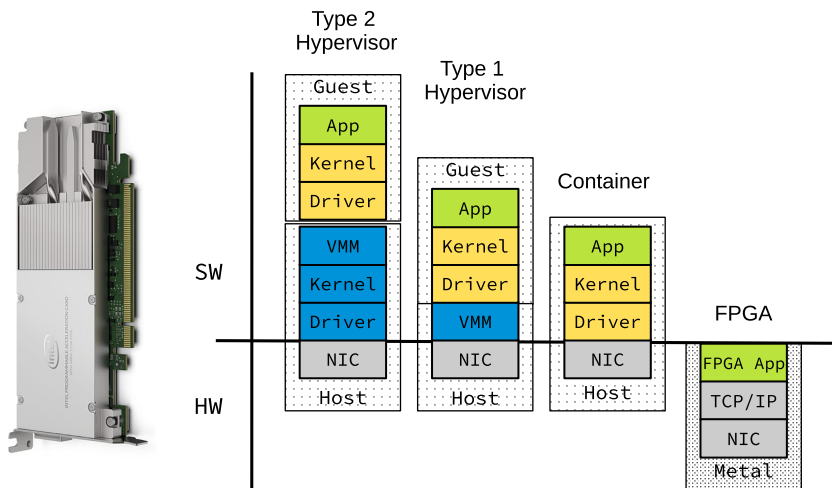
└ Introduction

└ Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.




TCP/IP in hardware using SME

Introduction

Background and Motivation

Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.



Applikationer og Big-Data udregninger flytter til Cloud, drevet af store data centre.



Disse data-centre kræver rigtigt meget plads, store mængder af strøm og er i stigende grad svære at vedligeholde og udvide.

De fleste data-centre er derfor begyndt at aflaste beregningerne til FPGAer, som fjerner meget af overhead til beregningerne


FPGA bruges til at få en computer til at køre hurtigere hvis de mest brugte instruktioner, skrives direkte ned i hardwaren

PROBLEMET er at der kun kan være en begrænset antal af FPGAer i konventionelle servere

A conventional data center architecture



Proposed disaggregated data center architecture (Weerasinghe et al. [2016])




TCP/IP in hardware using SME


└ Introduction

2019-09-18

Proposed disaggregated data center architecture (Weerasinghe et al. [2016])



Hvis man splitter ressourcerne op, kan man takket været FPGA få bedre ydeevne på det samme areal, samt nemmere håndtering af servere og deres komponenter.





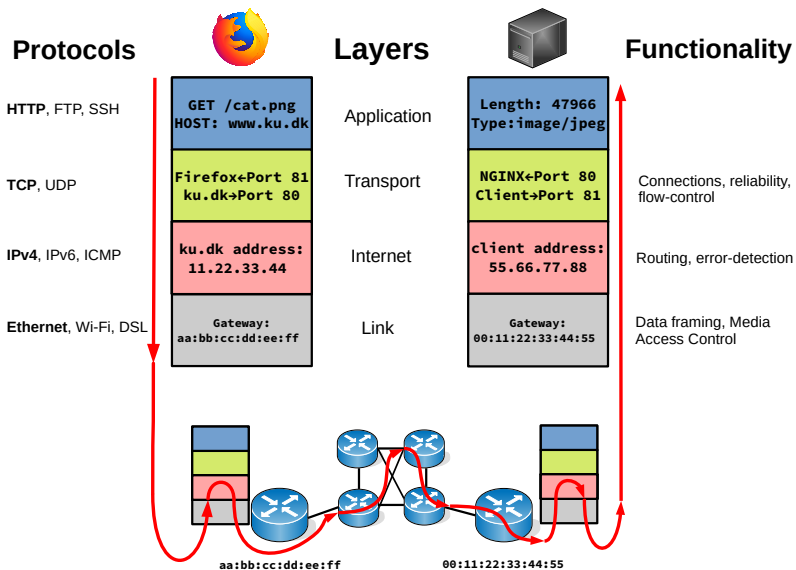


Figure: Map of about 30% of the accessible the endpoints on the Internet



The Internet Protocol Suite – A scenario




TCP/IP in hardware using SME


Introduction

2019-09-18

The Internet Protocol Suite – A scenario



Design with the 4 layers in mind



TCP/IP in hardware using SME

└ Introduction

2019-09-18

Design with the 4 layers in mind




Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME
└─ Implementation
└─ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

- SME introduction
- Processes
 - State machines
- Buffers
 - Memory segments
 - Dictionary
- Interface signal control
 - Buffer-Producer
 - Compute-Producer
- Interface control
 - Usage
 - Limitations

SME(Synchronous Message Exchange) introduction

- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing

TCP/IP in hardware using SME

└ Implementation

SME(Synchronous Message Exchange) introduction

- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing



2019-09-18

- a

State machines

```
1 public class SomeProcess :  
2   ↳ StateProcess  
3 {  
4   private override async  
5     ↳ Task OnTickAsync()  
6   {  
7     a();  
8     await ClockAsync();  
9     b();  
10    await ClockAsync();  
11    c();  
12    await ClockAsync();  
13  }
```

```
1 public class SomeProcess :  
2   ↳ SimpleProcess  
3 {  
4   // Initial state  
5   state = A;  
6  
7   protected override void  
8     ↳ OnTick()  
9   {  
10    switch(state) {  
11      case A:  
12        a();  
13        state = B;  
14      case B:  
15        b();  
16        state = C;  
17      case C:  
18        c();  
19        state = A;  
20    }  
21  }
```



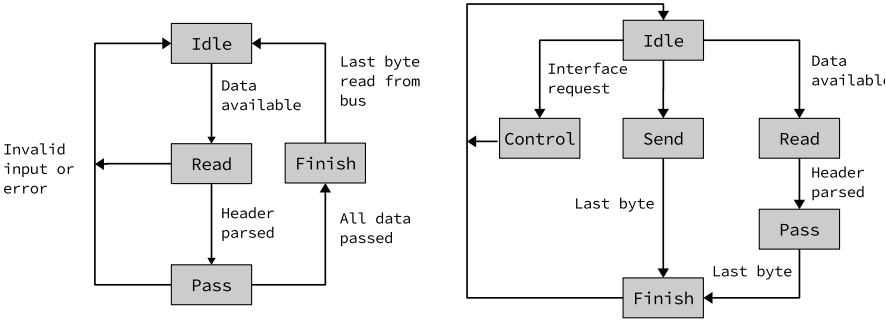
State machines

- **StateProcess**
Eksekvering kan stoppes når som helst(i bidder)
- **SimpleProcess**
Run er en clock altid, state machine håndteres med en switchcase.
Algoritme kan splittes op i flere bidder, men kræver en state per bid

Implementation

Processes

Examples




The internet process state machine The transport process state machine

2019-09-18

TCP/IP in hardware using SME

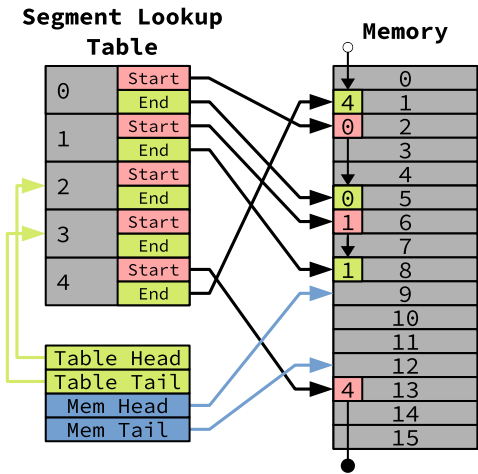
└ Implementation

└ Implementation



Labels!

Memory segments



TCP/IP in hardware using SME


└ Implementation

└ Implementation

Implementation



Buffers

Memory segments



Thumbnail?

Memory dictionary




Statisk tabel?

Animationer? Billede af splitup af pakker? Hold det uden protocol specifik info.

Implementation


Buffers

Some problems with the memory dictionaries!

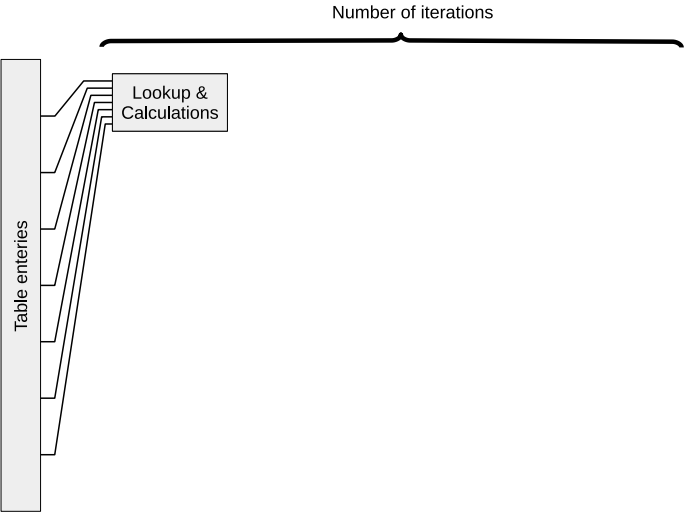


2019-09-18


- TCP/IP in hardware using SME
 - Implementation
 - Implementation



Some problems with the memory dictionaries!




- 2019-09-18 TCP/IP in hardware using SME
 - Implementation
 - Implementation



Implementation

Buffers


Some problems with the memory dictionaries!



- 2019-09-18 TCP/IP in hardware using SME
 - Implementation
 - Implementation

Implementation


Some problems with the memory dictionaries!



Implementation

Buffers

Some problems with the memory dictionaries!




TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation Buffers

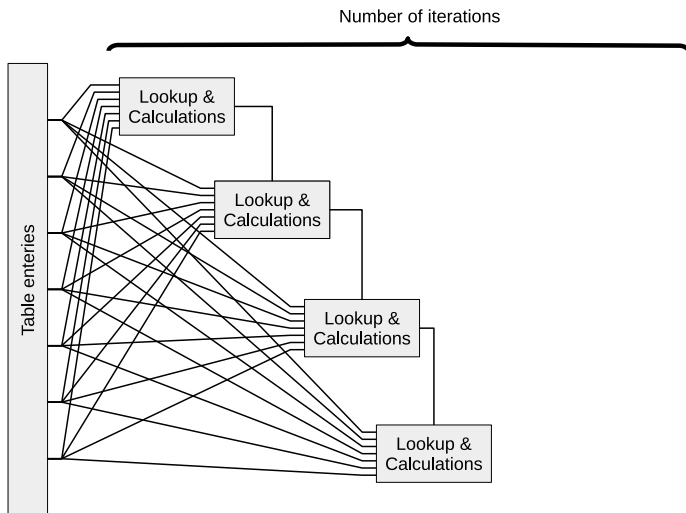
Some problems with the memory dictionaries!



Implementation

Buffers

Some problems with the memory dictionaries!




TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation Buffers


Some problems with the memory dictionaries!



Implementation

Buffers


Some problems with the memory dictionaries!



TCP/IP in hardware using SME

└ Implementation


└ Implementation



Implementation

Interface signal protocol

Identifying the scenarios




TCP/IP in hardware using SME

Implementation

Implementation

Implementation
Interface signal protocol
Identifying the scenarios



Data skal overføres hurtigst muligt, og det må ikke gå tabt


2 scenarier: fra "compute" til buffer, og omvendt

- CP kan ikke vente
- BP har stor buffer, og consumer starter transaktion

Implementation

Interface signal protocol

Compute-Producer (CP)




2019-09-18 TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation


Interface signal protocol



Implementation

Interface signal protocol

Buffer-Producer: Inspired by AXI4




TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation Interface signal protocol

Buffer-Producer: Inspired by AXI4




Implementation

Interface signal protocol

Buffer-Producer (BP)

Producer

Consumer



TCP/IP in hardware using SME

Implementation

Implementation

Implementation

Interface signal protocol




Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18
TCP/IP in hardware using SME
└─ Evaluation
└─ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 **Evaluation**
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

- Setup
 - Graph file simulator
- Test
- Validation
 - Latency
 - Outgoing packet validation
 - Internet Protocol Suite compliancy as per RFC 1122

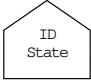
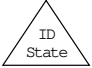
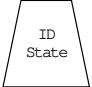

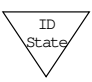

Graph file simulation







- Full input - output
- Does not take latency between packets into account
- Simplifies test cases

Definer send og receive bedre

Evaluation

Setup

Graph simulation node types			State and Color	Description
 Data In	 Send	 Command	Waiting	Vertex is not in use.
			Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
 Data Out	 Receive	 Wait	Inactive	Vertex is inactive. Simulator is not gathering data.
			Done	Vertex is done and validated.

Evaluation Setup			State and Color	Description
Graph simulation node types			Waiting	Vertex is not in use.
 Data In	 Send	 Command	Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
			Inactive	Vertex is inactive. Simulator is not gathering data.
 Data Out	 Receive	 Wait	Done	Vertex is done and validated.


Setup



└ Evaluation

└ Evaluation

2019-09-18



Senario

- Real life scenario
- Test at high workloads
- Remove garbage
- Respond to packet
- Differ between concurrent connections

The test

- 17283 packets in total
- Two "sessions"
- 640*2 UDP packets that needs a response
- 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

TCP/IP in hardware using SME

└─Evaluation

└─Evaluation

Evaluation Test

The test

- 17283 packets in total
 - Two "sessions"
 - 640*2 UDP packets that needs a response
 - 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

Latency calculations:

n_D : The number of bytes in the data part of the protocol. This excludes both headers from transport and internet.

n_I : The internet header size.

n_T : The transport header size.

n : The total packet size.

From packet to user

$$6 + n_I + 2n_T + 3n_D$$

From user to packet

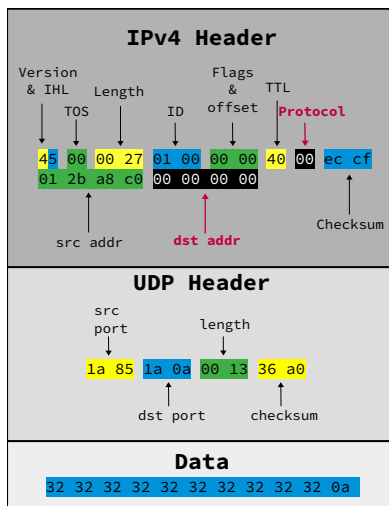
$$8 + 2n_I + 3n_T + 4n_D$$

Hav billede af sytem ved siden af

system graf med selve latency

bufferen kan ikke videresende data dirrekte, da den skal gemme segmentet først

Outgoing packet validation:



TCP/IP in hardware using SME


└ Evaluation

└ Evaluation

Evaluation

Validation

Outgoing packet validation:



Protocol ikke sat korrekt, destination ip ikke sat korrekt

Internet Protocol Suite compliancy as per RFC 1122

2019-09-18	TCP/IP in hardware using SME		Evaluation
	└─Evaluation		Validation
	└─Evaluation		Internet Protocol Suite compliancy as per RFC 1122

Ikke testet helt igennem, mem felterne er generelt sat

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion**
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME


└ Discussion

└ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion**
- 5 Conclusion
- 6 Future Work
- 7 Questions

Improving the performance:



Estimated performance:

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

TCP/IP in hardware using SME

└ Discussion

└ Discussion

Discussion

Estimated performance
 $1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

Improving the performance:




Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion**
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME
└─ Conclusion
└─ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion**
- 6 Future Work
- 7 Questions

- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

TCP/IP in hardware using SME

└ Conclusion

└ Conclusion

Conclusion

- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work**
- 7 Questions


2019-09-18 TCP/IP in hardware using SME
└ Future Work
└ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work**
- 7 Questions

Future Work

Firewall




TCP/IP in hardware using SME

└ Future Work


└ Future Work

Future Work
Firewall



Integration med buffere. Hvad ville det indebære

Implementing TCP



TCP/IP in hardware using SME

└Future Work

└Future Work

Future Work

TCP

Implementing TCP




Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions**

2019-09-18

TCP/IP in hardware using SME

└─ Questions

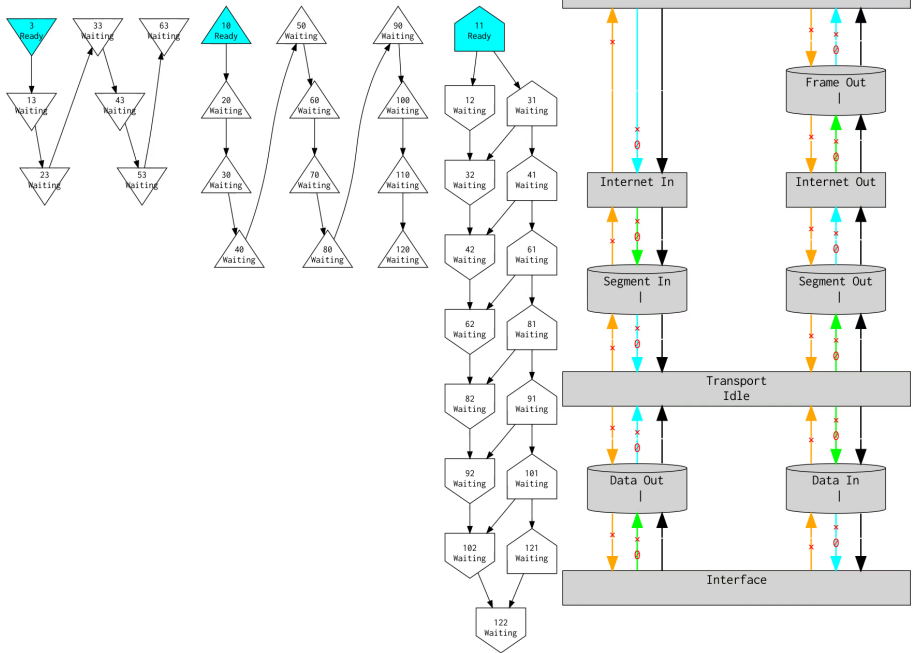
└─ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions**

- [1] J. Weerasinghe, F. Abel, C. Hagleitner, and A. Herkersdorf. Disaggregated fpgas: Network performance comparison against bare-metal servers, virtual machines and linux containers. In *2016 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*, pages 9–17, Dec 2016. doi: 10.1109/CloudCom.2016.0018.

clock: -1



TCP/IP in hardware using SME

Questions

