

# TCP/IP in hardware using SME

Mark Jan Jacobi & Jan Meznik

KU

September 18, 2019

## TCP/IP in hardware using SME

2019-09-18

**Mark** siger introduktion og 2-3 sætninger "abstrakt"

# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18

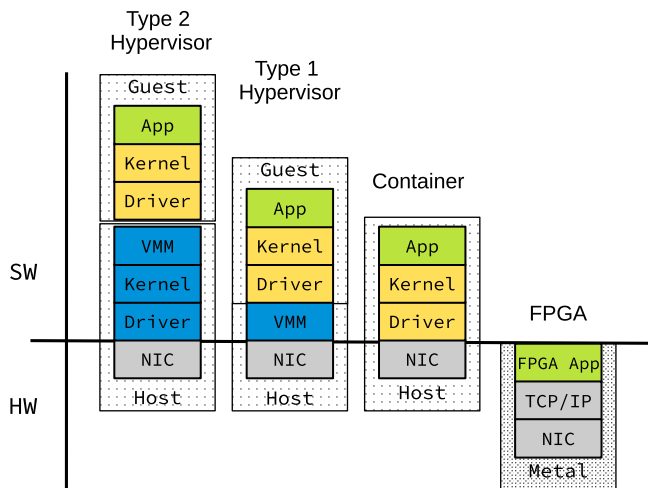
TCP/IP in hardware using SME

└ Introduction

└ Table of Contents

# Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.




## TCP/IP in hardware using SME

### └ Introduction

### └ Background and Motivation

#### Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.



Applikationer og Big-Data udregninger flytter til Cloud, drevet af store data centre.



Disse data-centre kræver rigtigt meget plads, store mængder af strøm og er i stigende grad svære at vedligeholde og udvide.

De fleste data-centre er derfor begyndt at aflaste beregningerne til FPGA'er, som fjerner meget af overhead til beregningerne


FPGA bruges til at få en computer til at køre hurtigere hvis de mest brugte instruktioner, skrives direkte ned i hardwaren

PROBLEMET er at der kun kan være en begrænset antal af FPGA'er i konventionelle servere

# A conventional data center architecture



# Proposed disaggregated data center architecture (Weerasinghe et al. [2016])




## TCP/IP in hardware using SME


### └ Introduction

2019-09-18

Proposed disaggregated data center architecture (Weerasinghe et al. [2016])



Hvis man splitter ressourcerne op, kan man takket været FPGA få bedre ydeevne på det samme areal, samt nemmere håndtering af servere og deres komponenter.





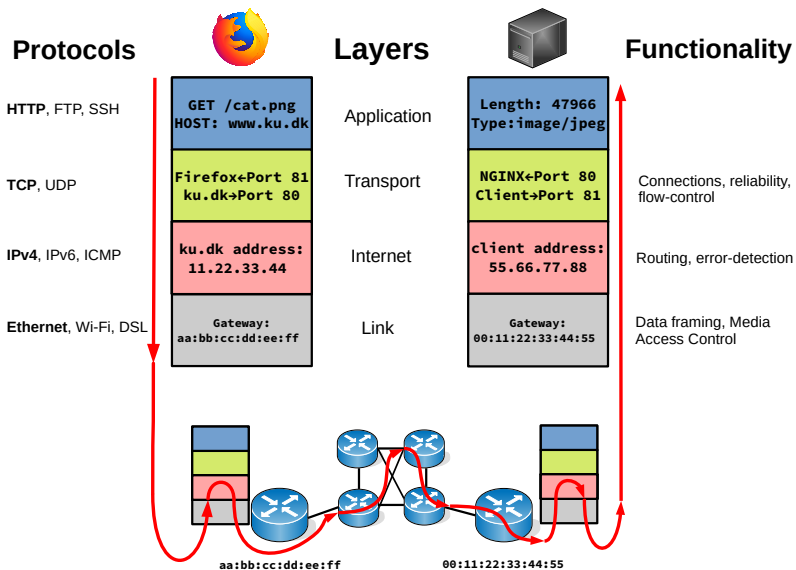


Figure: Map of about 30% of the accessible the endpoints on the Internet



# The Internet Protocol Suite – A scenario




## TCP/IP in hardware using SME

### Introduction


2019-09-18

The Internet Protocol Suite – A scenario





# Design with the 4 layers in mind



## TCP/IP in hardware using SME

### └ Introduction

2019-09-18

Design with the 4 layers in mind



# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME  
└ Implementation  
└ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

- SME introduction
- Processes
  - State machines
- Buffers
  - Memory segments
  - Dictionary
- Interface signal control
  - Buffer-Producer
  - Compute-Producer
- Interface control
  - Usage
  - Limitations

## SME(Synchronous Message Exchange) introduction

- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing

### TCP/IP in hardware using SME

#### └ Implementation

SME(Synchronous Message Exchange) introduction



- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing

- a

### State machines

```
1 public class SomeProcess :  
2   ↳ StateProcess  
3 {  
4   private override async  
5     ↳ Task OnTickAsync()  
6   {  
7     a();  
8     await ClockAsync();  
9     b();  
10    await ClockAsync();  
11    c();  
12    await ClockAsync();  
13  }
```

```
1 public class SomeProcess :  
2   ↳ SimpleProcess  
3 {  
4   // Initial state  
5   state = A;  
6   protected override void  
7     ↳ OnTick()  
8   {  
9     switch(state) {  
10      case A:  
11        a();  
12        state = B;  
13      case B:  
14        b();  
15        state = C;  
16      case C:  
17        c();  
18        state = A;  
19    }  
20  }
```



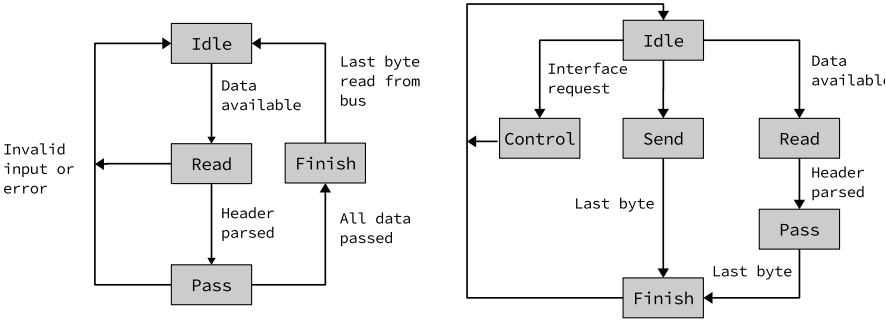
### State machines

- **StateProcess**  
Eksekvering kan stoppes når som helst(i bidder)
- **SimpleProcess**  
Run er en clock altid, state machine håndteres med en switchcase.  
Algoritme kan splittes op i flere bidder, men kræver en state per bid

# Implementation

## Processes

### Examples




The internet process state machine    The transport process state machine

2019-09-18

TCP/IP in hardware using SME

└ Implementation

└ Implementation

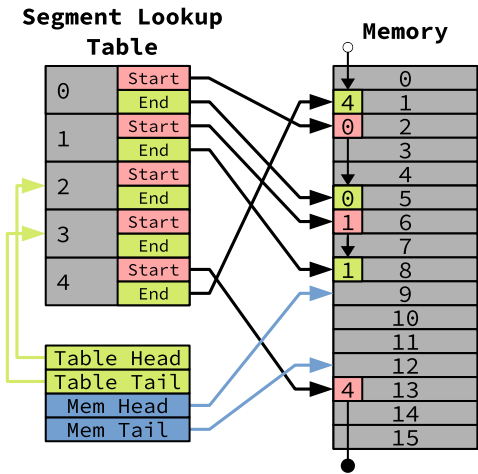


Labels!

# Implementation

## Buffers

### Memory segments



## TCP/IP in hardware using SME

└ Implementation


└ Implementation

Thumbnail?


Implementation

Buffers

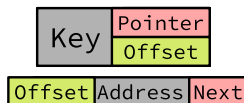
Memory segments



### Memory dictionary




### Value Legend



### TCP/IP in hardware using SME

└ Implementation

└ Implementation



Statisk tabel?


Animationer? Billede af splitup af pakker? Hold det uden protocol specifik info.



# Implementation


## Buffers

Some problems with the memory dictionaries!

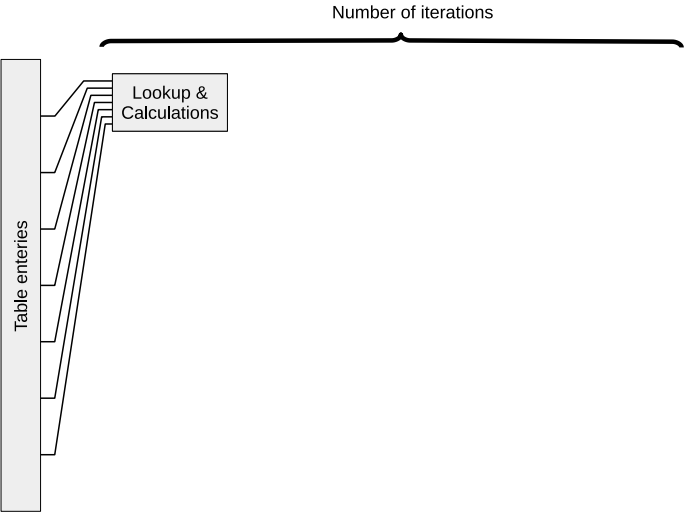


2019-09-18

- TCP/IP in hardware using SME
  - Implementation
  - Implementation



Some problems with the memory dictionaries!



## TCP/IP in hardware using SME


└ Implementation

└ Implementation

### Implementation

#### Buffers

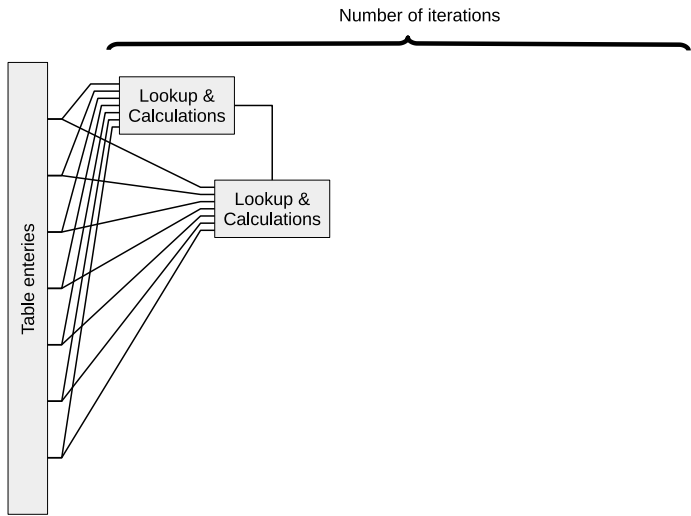
Some problems with the memory dictionaries!



# Implementation

## Buffers

Some problems with the memory dictionaries!




### TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation  
Buffers


Some problems with the memory dictionaries!



# Implementation

## Buffers


Some problems with the memory dictionaries!



### TCP/IP in hardware using SME

└ Implementation


└ Implementation



# Implementation

## Buffers

Some problems with the memory dictionaries!




## TCP/IP in hardware using SME

└ Implementation

└ Implementation

### Implementation Buffers


Some problems with the memory dictionaries!



# Implementation

## Buffers


Some problems with the memory dictionaries!



## TCP/IP in hardware using SME

└ Implementation


└ Implementation



# Implementation

## Interface signal protocol

### Identifying the scenarios




## TCP/IP in hardware using SME

### Implementation

### Implementation

Implementation  
Interface signal protocol


Identifying the scenarios



# Implementation

## Interface signal protocol

### Compute-Producer (CP)



## TCP/IP in hardware using SME

### Implementation

### Implementation

#### Implementation

interface signal protocol






# Implementation

## Interface signal protocol

### Buffer-Producer: Inspired by AXI4

- Single clock offset when sending data.
- Indicate end of stream with `bytes_left`.



## TCP/IP in hardware using SME


### └ Implementation

### └ Implementation

#### Implementation Interface signal protocol

##### Buffer-Producer: Inspired by AXI4

- Single clock offset when sending data.
- Indicate end of stream with `bytes_left`.




# Implementation

## Interface signal protocol

### Buffer-Producer (BP)

Producer

Consumer




## TCP/IP in hardware using SME

### Implementation

### Implementation

### Implementation

### Interface signal protocol



# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18

TCP/IP in hardware using SME

└─ Evaluation

└─ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

- Setup
  - Graph file simulator
- Test
- Validation
  - Latency
  - Outgoing packet validation
  - Internet Protocol Suite compliancy as per RFC 1122

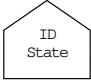
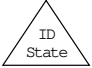
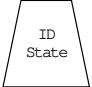

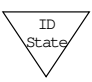

### Graph file simulation

- Full input - output
- Does not take latency between packets into account
- Simplifies test cases

Definer send og receive bedre

# Evaluation







## Setup

Graph simulation node types			State and Color	Description
 Data In	 Send	 Command	Waiting	Vertex is not in use.
			Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
 Data Out	 Receive	 Wait	Inactive	Vertex is inactive. Simulator is not gathering data.
			Done	Vertex is done and validated.

## TCP/IP in hardware using SME

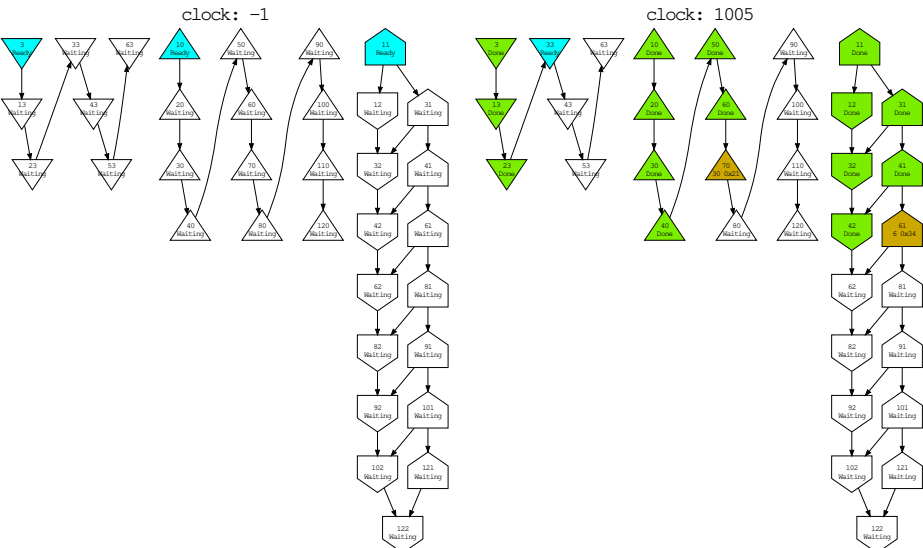
└ Evaluation

└ Evaluation

Evaluation Setup			State and Color	Description
Graph simulation node types			Waiting	Vertex is not in use.
 Data In	 Send	 Command	Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
			Inactive	Vertex is inactive. Simulator is not gathering data.
 Data Out	 Receive	 Wait	Done	Vertex is done and validated.

# Evaluation


## Setup



## TCP/IP in hardware using SME

└ Evaluation

└ Evaluation



### Senario

- Real life scenario
- Test at high workloads
- Remove garbage
- Respond to packet
- Differ between concurrent connections



### The test

- 17283 packets in total
- Two "sessions"
- 640\*2 UDP packets that needs a response
- 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

Latency calculations:

$n_D$  : The number of bytes in the data part of the protocol. This excludes both headers from transport and internet.

$n_I$  : The internet header size.

$n_T$  : The transport header size.

$n$  : The total packet size.

From packet to user

$$6 + n_I + 2n_T + 3n_D$$

From user to packet

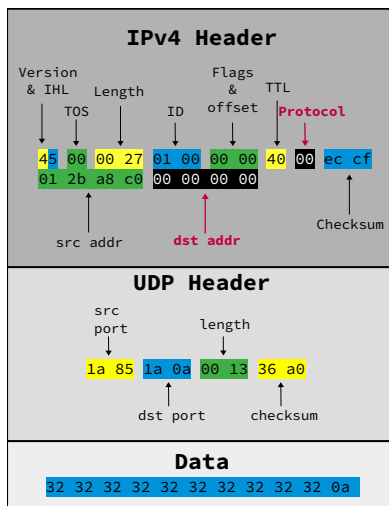
$$8 + 2n_I + 3n_T + 4n_D$$

Hav billede af sytem ved siden af

system graf med selve latency

bufferen kan ikke videresende data dirrekte, da den skal gemme segmentet først

Outgoing packet validation:



TCP/IP in hardware using SME


└ Evaluation

└ Evaluation

Evaluation

Validation

Outgoing packet validation:



Protocol ikke sat korrekt, destination ip ikke sat korrekt

## Internet Protocol Suite compliancy as per RFC 1122

2019-09-18

TCP/IP in hardware using SME	
└─Evaluation	
└─Evaluation	

Ikke testet helt igennem, mem felterne er generelt sat

# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion**
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME


└ Discussion

└ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion**
- 5 Conclusion
- 6 Future Work
- 7 Questions

Improving the performance:



Estimated performance:

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$


TCP/IP in hardware using SME

└ Discussion

└ Discussion

Discussion

Estimated performance  
 $1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$



Usability

SOMETHING

2019-09-18

TCP/IP in hardware using SME

└ Discussion

Using C#

State modelling  
Simulation  
Concurrency

2019-09-18

TCP/IP in hardware using SME

└ Discussion

Using C#

State modelling  
Simulation  
Concurrency



# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion**
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME

└ Conclusion

└ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion**
- 6 Future Work
- 7 Questions

- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

## TCP/IP in hardware using SME

### └ Conclusion

### └ Conclusion

- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work**
- 7 Questions

2019-09-18


- TCP/IP in hardware using SME
  - └ Future Work
    - └ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work**
- 7 Questions

# Future Work

## Firewall




## TCP/IP in hardware using SME

└ Future Work

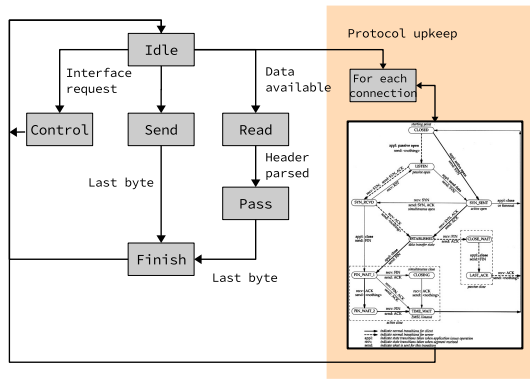
└ Future Work

Future Work  
Firewall



Integration med buffere. Hvad ville det indebære

Implementing TCP



TCP/IP in hardware using SME


└Future Work

└Future Work

Future Work

TCP

Implementing TCP



# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions**

2019-09-18

TCP/IP in hardware using SME

└─ Questions

└─ Table of Contents

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions**

- [1] J. Weerasinghe, F. Abel, C. Hagleitner, and A. Herkersdorf. Disaggregated fpgas: Network performance comparison against bare-metal servers, virtual machines and linux containers. In *2016 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*, pages 9–17, Dec 2016. doi: 10.1109/CloudCom.2016.0018.

