

TCP/IP in hardware using SME

Mark Jan Jacobi & Jan Meznik

KU

September 19, 2019

TCP/IP in hardware using SME

2019-09-19

Mark siger introduktion og 2-3 sætninger "abstrakt"

Table of Contents

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion

6 Future Work

7 Questions

8 Demonstration

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion


6 Future Work

7 Questions

8 Demonstration

Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.




TCP/IP in hardware using SME

└ Introduction

└ Background and Motivation

Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.



Applikationer og Big-Data udregninger flytter til Cloud, drevet af store data centre.


Disse data-centre kræver rigtigt meget plads, store mængder af strøm og er i stigende grad svære at vedligeholde og udvide.

De fleste data-centre er derfor begyndt at aflaste beregningerne til FPGA'er, som fjerner meget af overhead til beregningerne


FPGA bruges til at få en computer til at køre hurtigere hvis de mest brugte instruktioner, skrives direkte ned i hardwaren

PROBLEMET er at der kun kan være en begrænset antal af FPGA'er i konventionelle servere

A conventional data center architecture



Proposed disaggregated data center architecture (Weerasinghe et al. [2016])




TCP/IP in hardware using SME


Introduction

2019-09-19

Proposed disaggregated data center architecture (Weerasinghe et al. [2016])



Hvis man splitter ressourcerne op, kan man takket været FPGA få bedre ydeevne på det samme areal, samt nemmere håndtering af servere og deres komponenter.





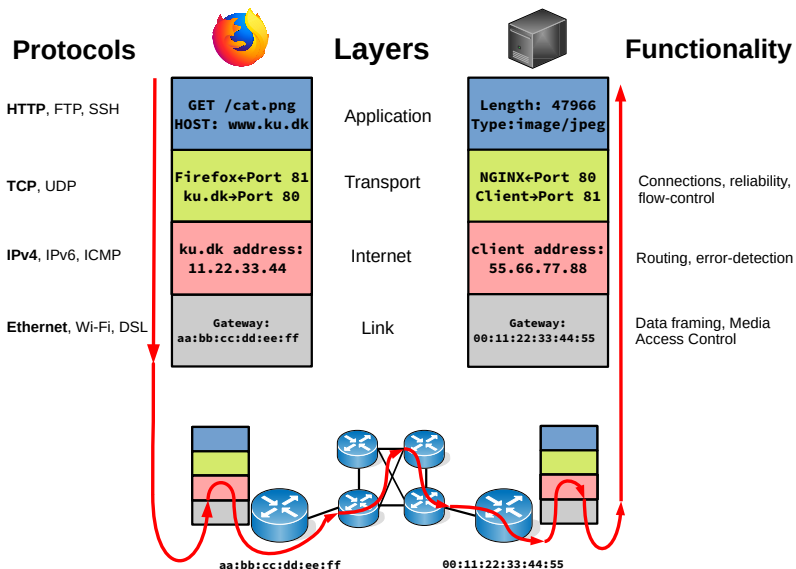


Figure: Map of about 30% of the accessible the endpoints on the Internet



The Internet Protocol Suite – A scenario




TCP/IP in hardware using SME


Introduction

2019-09-19

The Internet Protocol Suite – A scenario



Design with the 4 layers in mind




TCP/IP in hardware using SME

└ Introduction

2019-09-19

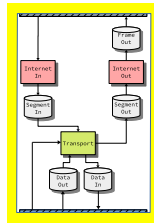
Design with the 4 layers in mind



- 1 Introduction
- 2 **Implementation**
 - SME introduction
 - Processes
 - Buffers
 - Interface signal protocol
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions
- 8 Demonstration

Implementation

SME introduction



SME(Synchronous Message Exchange) introduction

- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing
- Not fully feature complete with C#(No threads, no allocation)

TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation SME introduction


SME(Synchronous Message Exchange) introduction

- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing
- Not fully feature complete with C#(No threads, no allocation)

- What is a bus and a process
- No VHDL code
- Clocks abstracted away behind the management of processes and busses
- Testing straight in the simulator, but also in afterwards in the GHDL compiler, via an clock lookup table
- Since not feature complete, only simple structures can be used. We choose state diagrams since they are possible to make, and easy to understand

Processes

State machines



```


1 public class SomeProcess :
2     ↳ StateProcess
3 {
4     private override async
5         ↳ Task
6         ↳ OnTickAsync()
7 {
8     a();
9     await ClockAsync();
10    b();
11    await ClockAsync();
12    c();
13    await ClockAsync();
14 }
15 }

```

```

1 public class SomeProcess :
2     ↳ SimpleProcess
3 {
4     // Initial state
5     state = A;
6
7     protected override void
8         ↳ OnTick()
9     {
10         switch(state) {
11             case A:
12                 a();
13                 state = B;
14             case B:
15                 b();
16                 state = C;
17             case C:
18                 c();
19                 state = A;
20         }
21     }
22 }

```




TCP/IP in hardware using SME

Implementation

└ Implementation

2019-09-19

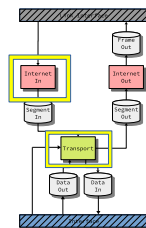


State machines


- `StateProcess`
Eksekvering kan stoppes når som helst(i bidder)
- `SimpleProcess`
Run er en clock altid, state machine håndteres med en switchcase.
Algoritme kan splittes op i flere bidder, men kræver en state per bid

Implementation

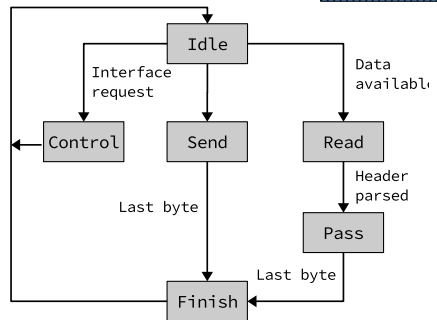
Processes



Examples



Internet in process state machine



Transport process state machine


TCP/IP in hardware using SME

Implementation

Implementation

Implementation Processes


Examples



- Gå igennem state diagrammer
- Snak om grundlaget for de forskellige typer brug

Implementation

Buffers



Why buffers?

- Fixes segmentation
- Processes can get data at their leisure

TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation Buffers

Why buffers?


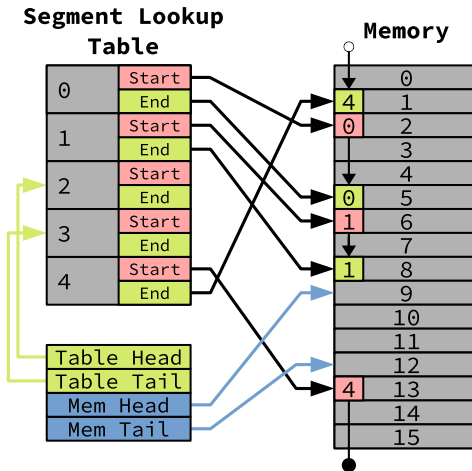
- Fixes segmentation
- Processes can get data at their leisure

Hvorfor bruer vi buffers?

Implementation

Buffers

Memory segments



TCP/IP in hardware using SME

Implementation

Implementation


Implementation
Buffers
Memory segments



- Reason behind?
 - Segment handling
 - References to other segment to concatting of segments later


Implementation

Buffers



Memory dictionary

Key Table Value Table



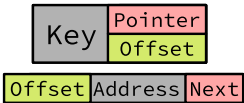
Initial state.
Key 5 have 2
elements:

[4, 18]

at index:

[2, 7]

Value Legend



TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation

Memory dictionary

Initial state.
Key 5 have 2
elements:

[4, 18]

at index:

[2, 7]

Value Legend

Key	Pointer
	Offset

Offset	Address	Next

Snak om input


2019-09-19

Implementation

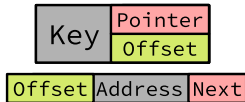
Buffers

Memory dictionary

Key Table Value Table



Value Legend




Insert element 8: Key 5 have 3 elements:

[4, 8, 18]

at index:

[2, 6, 7]




TCP/IP in hardware using SME

Implementation

Implementation

Snak om input




Implementation

Buffers

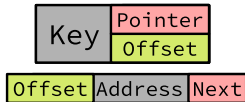
Memory dictionary

Key Table

Value Table



Value Legend




Insert element 2: Key 5 have 4 elements:

[2, 4, 8, 18]

at index:

[5, 2, 6, 7]



TCP/IP in hardware using SME

Implementation

Implementation

Snak om input

Implementation

Memory dictionary

Insert element 2: Key 5 have 4 elements: [2, 4, 8, 18] at index: [5, 2, 6, 7]

Key Table

Key	Pointer
0	0
1	1
5	5
-1	-1

Value Table

Value	Offset
2	0
5	1
4	2
2	3
0	4
2	5
10	6
0	7

Value Legend


Key	Pointer
	Offset

Offset	Address	Next

Implementation


Buffers

Some problems with the memory dictionaries!



Number of iterations

2019-09-19 TCP/IP in hardware using SME
└ Implementation
└ Implementation





Overflow!
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation

Buffers

Some problems with the memory dictionary!

Overflow!


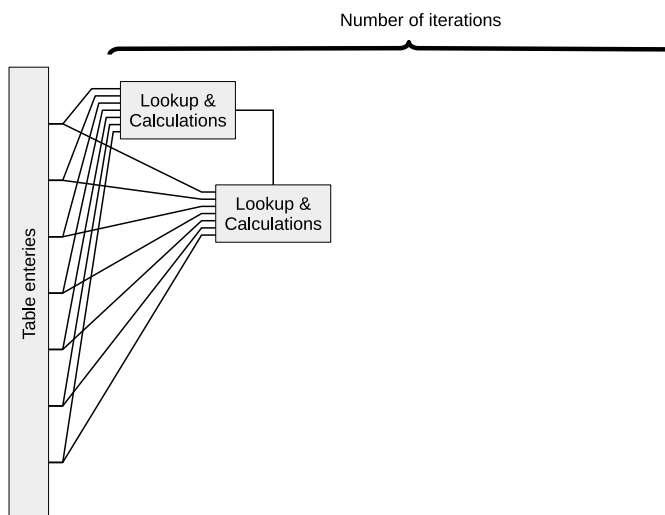
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME


Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!


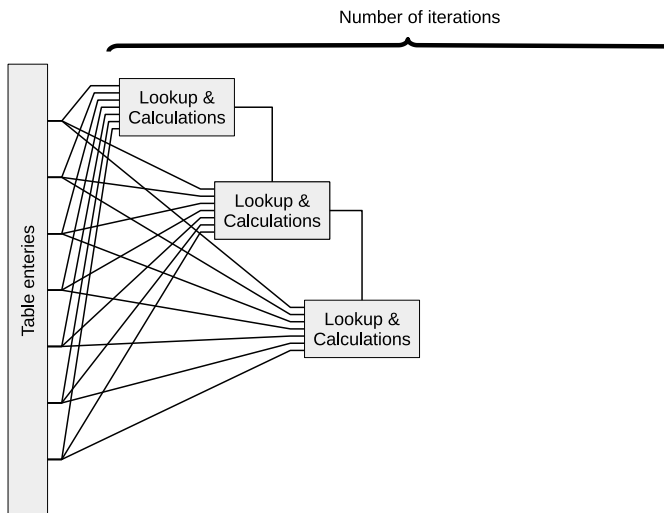
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME


Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!


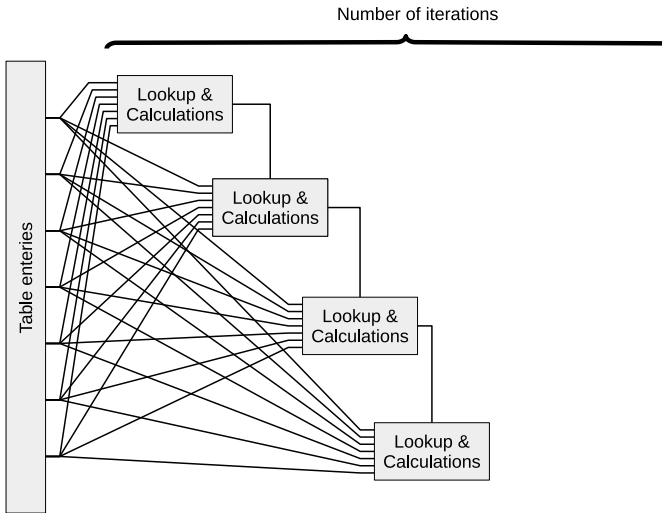
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME

Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!


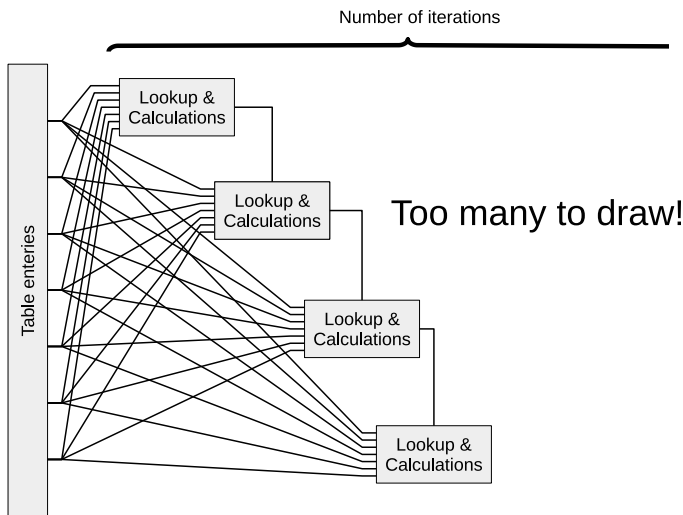
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME


Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!



kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Interface signal protocol

Identifying the scenarios



TCP/IP in hardware using SME

└ Implementation

└ Implementation


Implementation
Interface signal protocol
Identifying the scenarios




2019-09-19

Implementation

Interface signal protocol



Compute-Producer (CP)




2019-09-19 TCP/IP in hardware using SME

└ Implementation


└ Implementation

Implementation
Interface signal protocol




Implementation

Interface signal protocol



Buffer-Producer: Inspired by AXI4

- Single clock offset when sending data.
- Indicate end of stream with `bytes_left`.



TCP/IP in hardware using SME

Implementation


Implementation

Implementation

Interface signal protocol

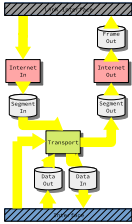
Buffer-Producer: Inspired by AXI4

- Single clock offset when sending data.
- Indicate end of stream with `bytes_left`.




Implementation

Interface signal protocol



Buffer-Producer (BP)

Producer Consumer



2019-09-19 TCP/IP in hardware using SME
└ Implementation

└ Implementation

Implementation
Interface signal protocol



1

Introduction

2

Implementation

3

Evaluation

• Setup

• Test

• Validation

4

Discussion

5

Conclusion

6

Future Work

7

Questions

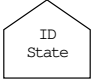
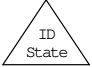
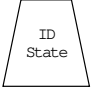



8

Demonstration

Graph file simulation

- Full input - output
- Does not take latency between packets into account
- Simplifies test cases

Definer send og receive bedre

Graph simulation node types			State and Color	Description
 Data In	 Send	 Command	Waiting	Vertex is not in use.
			Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
 Data Out	 Receive	 Wait	Inactive	Vertex is inactive. Simulator is not gathering data.
			Done	Vertex is done and validated.




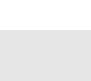
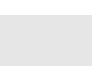

2019-09-19

TCP/IP in hardware using SME

└─Evaluation

└─Evaluation

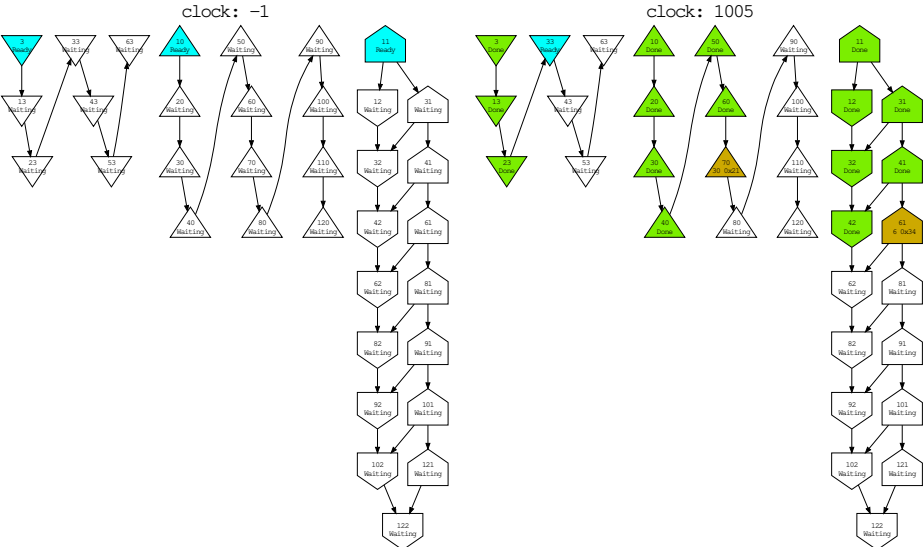
Evaluation
Setup

Graph simulation node types			State and Color	Description
 Data In	 Send	 Command	Waiting	Vertex is not in use.
			Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
 Data Out	 Receive	 Wait	Inactive	Vertex is inactive. Simulator is not gathering data.
			Done	Vertex is done and validated.

Hop til illustrationen på næste slide nå du snakker om det!

Evaluation

Setup



Mark Jan Jacobi & Jan Meznik (KU)


TCP/IP in hardware using SME

September 19, 2019 26 / 45

TCP/IP in hardware using SME

└ Evaluation

└ Evaluation



2019-09-19

Senario

- Real life scenario
- Test at high workloads
- Remove garbage
- Respond to packet
- Differ between concurrent connections

Fortæl kun om hvad vi vil have, ikke hvad vi har lavet af test

The test

- 17283 packets in total
- Two "sessions"
- 640*2 UDP packets that needs a response
- 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

- 17283 packets in total
 - Two "sessions"
 - 640*2 UDP packets that needs a response
 - 640 well formed UDP packets with no session (discard)
 - Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

Latency calculations:

n_D : The number of bytes in the data part of the protocol. This excludes both headers from transport and internet.

n_I : The internet header size.

n_T : The transport header size.


n : The total packet size.

From packet to user

$$6 + n_I + 2n_T + 3n_D$$

From user to packet

$$8 + 2n_I + 3n_T + 4n_D$$



TCP/IP in hardware using SME

└ Evaluation

└ Evaluation

Evaluation

Latency

Latency calculations:

n_D : The number of bytes in the data part of the protocol. This excludes both headers from transport and internet.

n_I : The internet header size.

n_T : The transport header size.


n : The total packet size.

From packet to user

$$6 + n_I + 2n_T + 3n_D$$


From user to packet

$$8 + 2n_I + 3n_T + 4n_D$$



Bufferen kan ikke videresende data direkte, da den skal gemme segmentet først

Outgoing packet validation:



TCP/IP in hardware using SME


└ Evaluation

└ Evaluation

Evaluation

Validation

Outgoing packet validation:



Protocol ikke sat korrekt, destination ip ikke sat korrekt

Internet Protocol Suite compliancy as per RFC 1122

2019-09-19


TCP/IP in hardware using SME	
└─Evaluation	
└─Evaluation	

Ikke testet helt igennem, mem felterne er generelt sat

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions
- 8 Demonstration

Improving the performance:



Estimated performance:

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

Discussion

Improving the performance:

Estimated performance

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

Usability

SOMETHING

2019-09-19

TCP/IP in hardware using SME

└ Discussion

Using C#

State modelling
Simulation
Concurrency

2019-09-19 TCP/IP in hardware using SME
└ Discussion

Using C#

State modelling
Simulation
Concurrency

Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions
- 8 Demonstration

- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

TCP/IP in hardware using SME

└ Conclusion

└ Conclusion

- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

Table of Contents

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion

6 Future Work

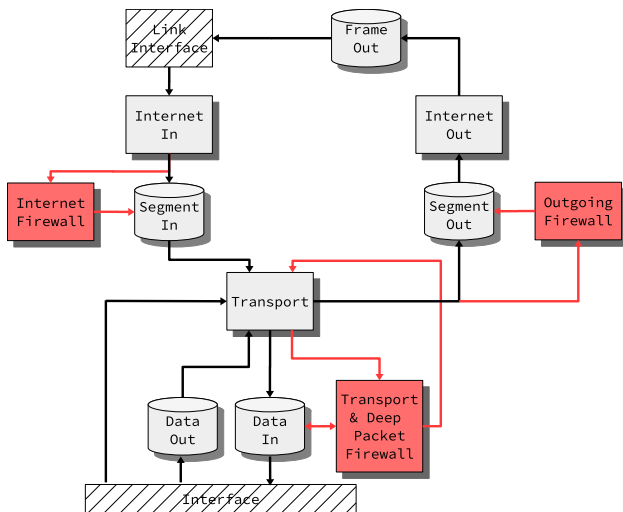
- Firewall
- TCP

7 Questions

8 Demonstration

Future Work

Firewall




TCP/IP in hardware using SME

└ Future Work

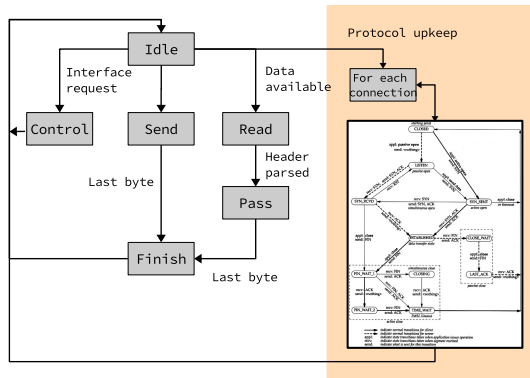
└ Future Work

Future Work
Firewall



Integration med buffere. Hvad ville det indebære

Implementing TCP



TCP/IP in hardware using SME

└ Future Work

└ Future Work

Future Work
TCP

Implementing TCP




Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions
- 8 Demonstration

- [1] J. Weerasinghe, F. Abel, C. Hagleitner, and A. Herkersdorf. Disaggregated fpgas: Network performance comparison against bare-metal servers, virtual machines and linux containers. In *2016 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*, pages 9–17, Dec 2016. doi: 10.1109/CloudCom.2016.0018.

- 1

Introduction
- 2

Implementation
- 3

Evaluation
- 4

Discussion
- 5

Conclusion
- 6

Future Work
- 7

Questions
- 8

Demonstration

Demonstration

TCP/IP in hardware using SME

└ Demonstration

└ Demonstration

Demonstration

