

TCP/IP in hardware using SME

Mark Jan Jacobi & Jan Meznik

KU

September 19, 2019

TCP/IP in hardware using SME

2019-09-19

TCP/IP in hardware using SME

Mark Jan Jacobi & Jan Meznik

KU

September 19, 2019

Mark siger introduktion og 2-3 sætninger "abstrakt"

Table of Contents

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion

6 Future Work

7 Questions

8 Demonstration

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion

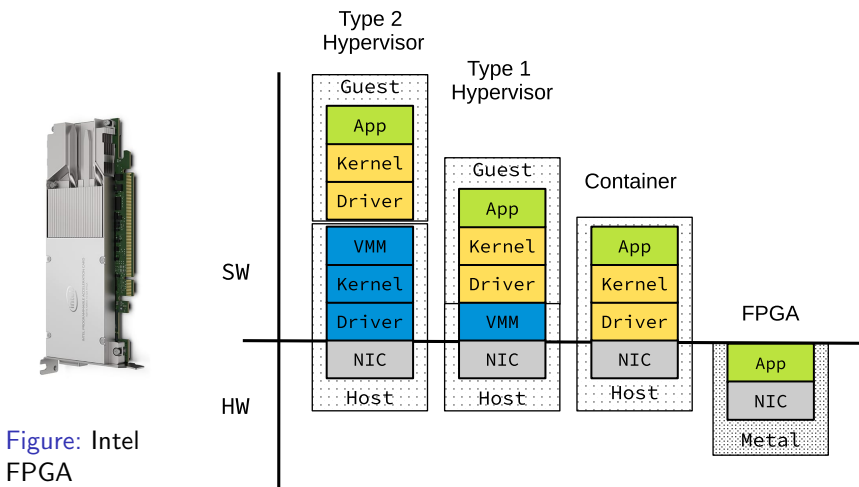
6 Future Work

7 Questions

8 Demonstration

Background and Motivation

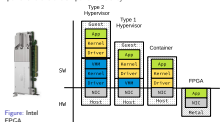
FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.



TCP/IP in hardware using SME

Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.



2019-09-19

Introduction

Background and Motivation

Applikationer og Big-Data udregninger flytter til Cloud, drevet af store data centre.

data-centre kraever meget plads, store maengder af stroem og er svaere at vedligeholde og udvide.

DC optimerer servere for at fa mest vaerdi muligt

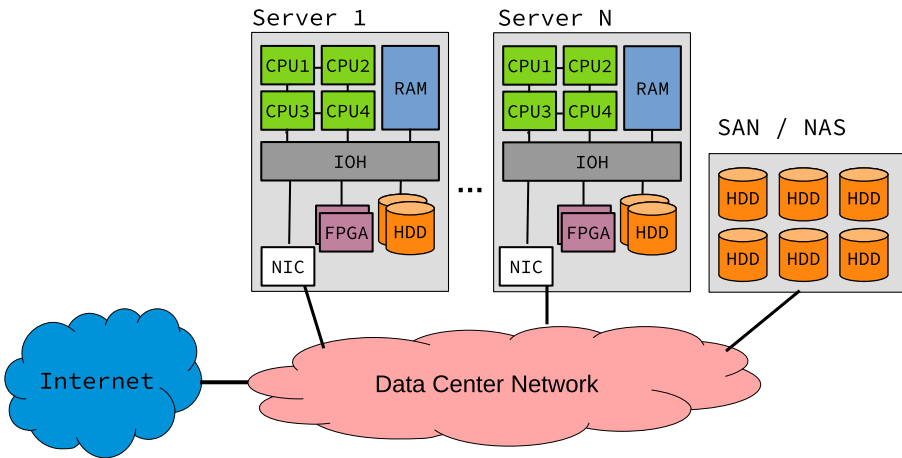
Tendens til aflaste beregninger til FPGAer, fjerne overhead

FPGA er hardware kan udføre beregninger hurtigt pga. dens parallelle programmerbare natur. Den er hurtigt fordi instruktioner skrives direkte ned i hardwaren

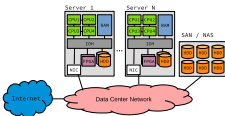
— GRAF HER —

PROBLEMET er at der kun kan vaere en begreanset antal af FPGAer i

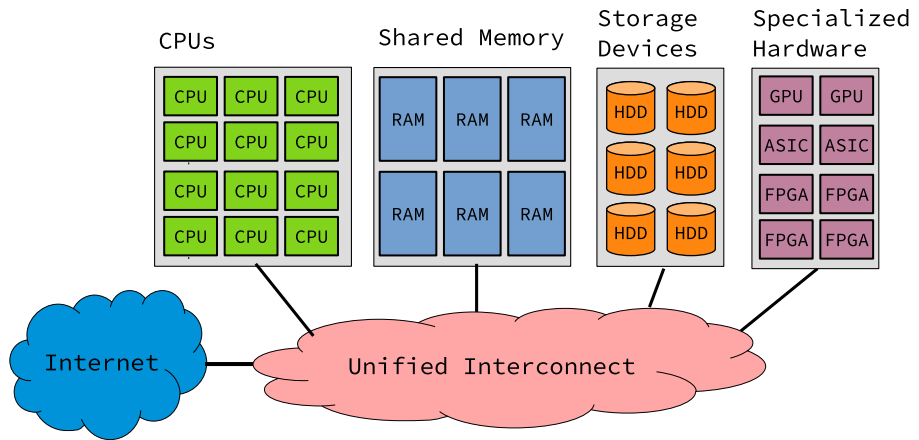
A conventional data center architecture



A conventional data center architecture



Proposed disaggregated data center architecture (Weerasinghe et al. [2016])

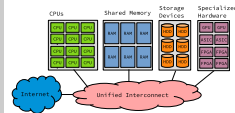


TCP/IP in hardware using SME

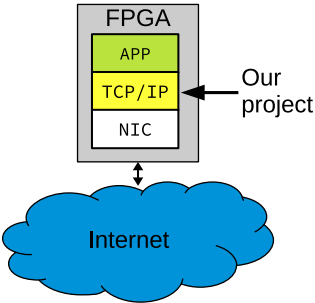
└ Introduction

2019-09-19

Proposed disaggregated data center architecture (Weerasinghe et al. [2016])



Hvis man splitter ressourcerne op, kan man takket været FPGA få bedre ydeevne på det samme areal, samt nemmere håndtering af servere og deres komponenter.



2019-09-19 TCP/IP in hardware using SME

└ Introduction

FPGA usage



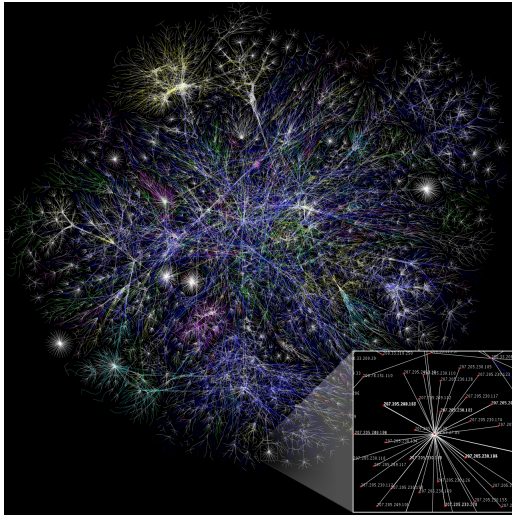
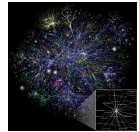
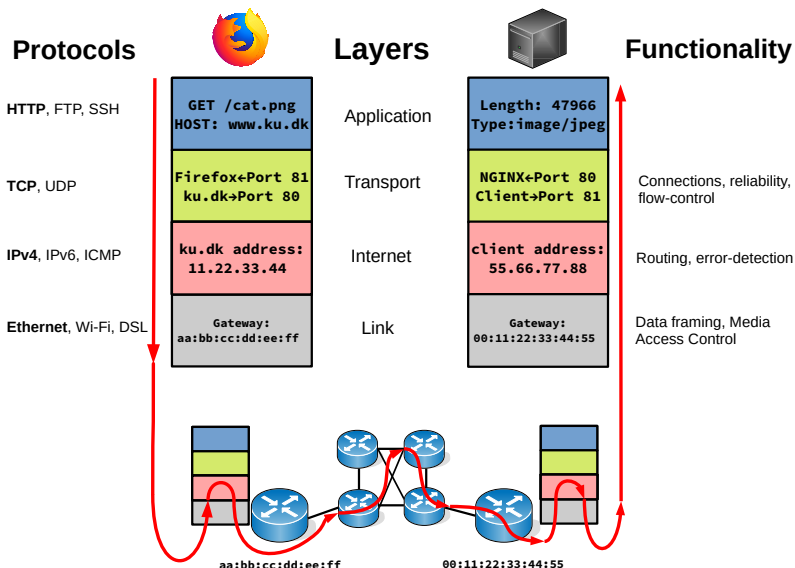


Figure: Map of about 30% of the accessible the endpoints on the Internet

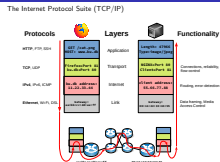


The Internet Protocol Suite (TCP/IP)



TCP/IP in hardware using SME

Introduction



TCP/IP er samling af standarder og protokoller

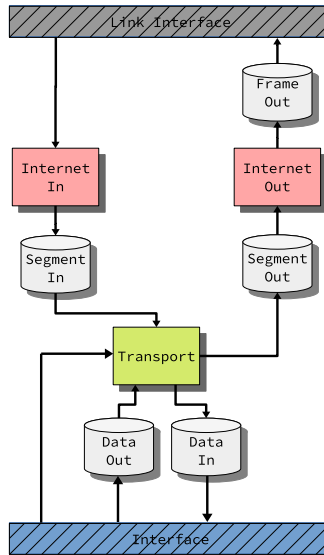
Link: Overførsel på det fysiske medium

Internet: bestemmer data-veje, adressering, fejl-kontrol

Transport: pålidelighed, forbindelser, kontrol flow

Application: Defineret af selve applikationen

Design with the 4 layers in mind



TCP/IP in hardware using SME

└ Introduction

2019-09-19

Design with the 4 layers in mind

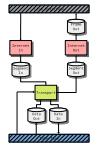


Table of Contents

1 Introduction

2 Implementation

- SME introduction
- Processes
- Buffers
- Interface signal protocol

3 Evaluation

4 Discussion

5 Conclusion

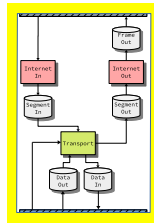
6 Future Work

7 Questions

8 Demonstration

Implementation

SME introduction



SME(Synchronous Message Exchange) introduction

- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing
- Not fully feature complete with C#(No threads, no allocation)

TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation SME introduction

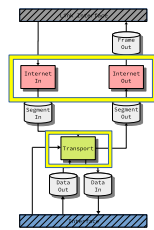
- SME(Synchronous Message Exchange) introduction
- Processes and Busses
 - Higher abstraction
 - Handling of clocks
 - Easy testing
 - Not fully feature complete with C#(No threads, no allocation)

- What is a bus and a process
- No VHDL code
- Clocks abstracted away behind the management of processes and busses
- Testing straight in the simulator, but also in afterwards in the GHDL compiler, via an clock lookup table
- Since not feature complete, only simple structures can be used. We choose state diagrams since they are possible to make, and easy to understand

Implementation

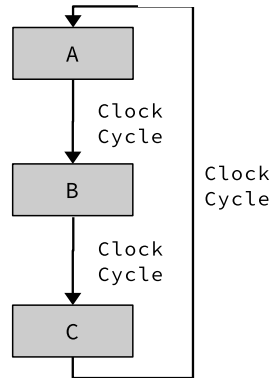
Processes

State machines



```
1 public class SomeProcess :  
2   ↳ StateProcess  
3 {  
4   private override async  
5     ↳ Task OnTickAsync()  
6   {  
7     a();  
8     await ClockAsync();  
9     b();  
10    await ClockAsync();  
11    c();  
12    await ClockAsync();  
13  }  
14 }
```

```
1 public class SomeProcess :  
2   ↳ SimpleProcess  
3 {  
4   // Initial state  
5   state = A;  
6  
7   protected override void  
8     ↳ OnTick()  
9   {  
10    switch(state) {  
11      case A: a(); state = B;  
12      case B: b(); state = C;  
13      case C: c(); state = A;  
14    }  
15  }  
16 }
```



TCP/IP in hardware using SME

Implementation

Implementation

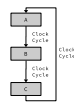
Implementation

Processes

State machines

```
1 public class SomeProcess :  
2   ↳ SimpleProcess  
3 {  
4   // Initial state  
5   state = A;  
6  
7   protected override void  
8     ↳ OnTick()  
9   {  
10    switch(state) {  
11      case A: a(); state = B;  
12      case B: b(); state = C;  
13      case C: c(); state = A;  
14    }  
15  }  
16 }
```

```
1 public class SomeProcess :  
2   ↳ SimpleProcess  
3 {  
4   // Initial state  
5   state = A;  
6  
7   protected override void  
8     ↳ OnTick()  
9   {  
10    switch(state) {  
11      case A: a(); state = B;  
12      case B: b(); state = C;  
13      case C: c(); state = A;  
14    }  
15  }  
16 }
```

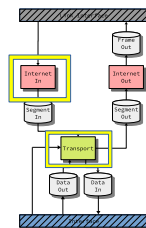


State machines

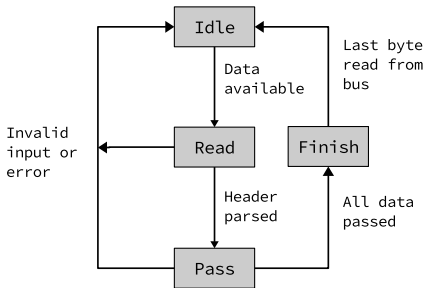
- **StateProcess**
Eksekvering kan stoppes når som helst(i bidder)
- **SimpleProcess**
Run er en clock altid, state machine håndteres med en switchcase.
Algoritme kan splittes op i flere bidder, men kræver en state per bid

Implementation

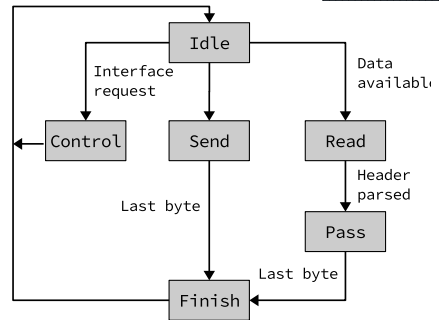
Processes



Examples



Internet in process state machine



Transport process state machine

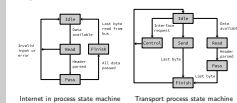
TCP/IP in hardware using SME

Implementation

Implementation

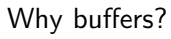
Implementation

Examples



- Gå igennem state diagrammer
- Snak om grundlaget for de forskellige typer brug

Buffers



- Mark Jan Jacobi & Jan Meznik (KU)

TCP/IP in hardware using SME

September 19, 2019

15 / 43

TCP/IP in hardware using SME

Implementation

Implementation

Implementation

Why buffers?

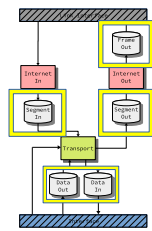
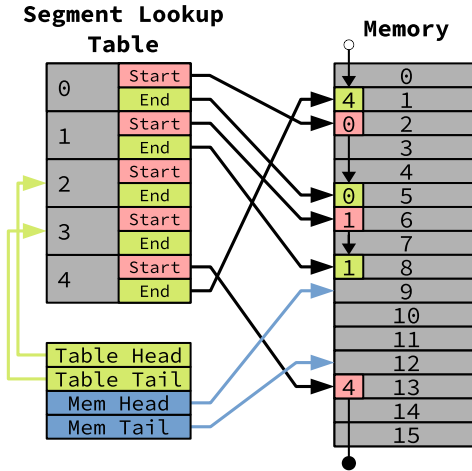
- ◆ Fixes segmentation
- ◆ Processes can get data at their leisure

Hvorfor bruer vi buffers?

Implementation

Buffers

Memory segments



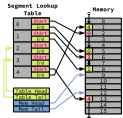
TCP/IP in hardware using SME

Implementation

Implementation

Implementation
Buffers

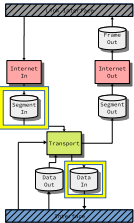
Memory segments



- Reason behind?
 - Segment handling
 - References to other segment to concatting of segments later

Implementation

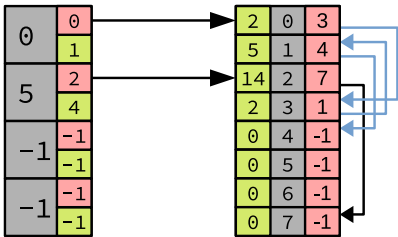
Buffers



Memory dictionary

Key Table

Value Table



Initial state.

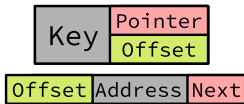
Key 5 have 2 elements:

[4, 18]

at index:

[2, 7]

Value Legend



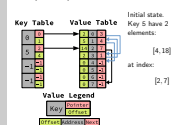
TCP/IP in hardware using SME

└ Implementation

└ Implementation

Implementation

Memory dictionary



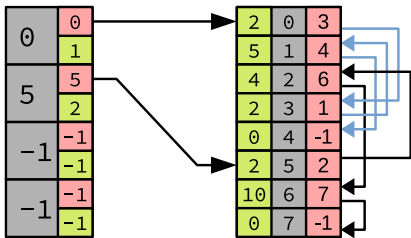
Snak om input

Implementation

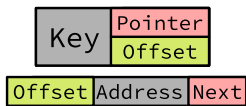
Buffers

Memory dictionary

Key Table Value Table



Value Legend

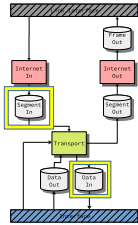


Insert element 2: Key 5 have 4 elements:

[2, 4, 8, 18]

at index:

[5, 2, 6, 7]



TCP/IP in hardware using SME

Implementation

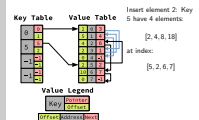
Implementation

Snak om input

Implementation

Buffers

Memory dictionary



Buffers

Table enteries

Implementation

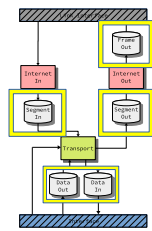
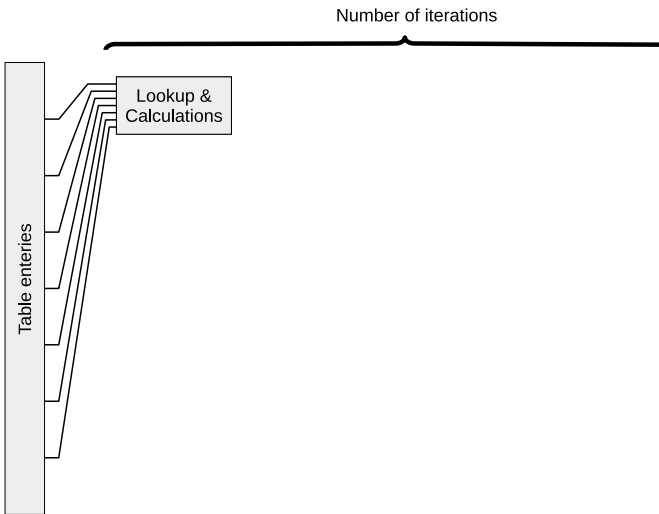
Some problems with the memory dictionaries!

kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME

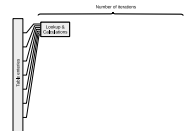
Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!

kan læses ved:

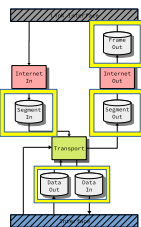
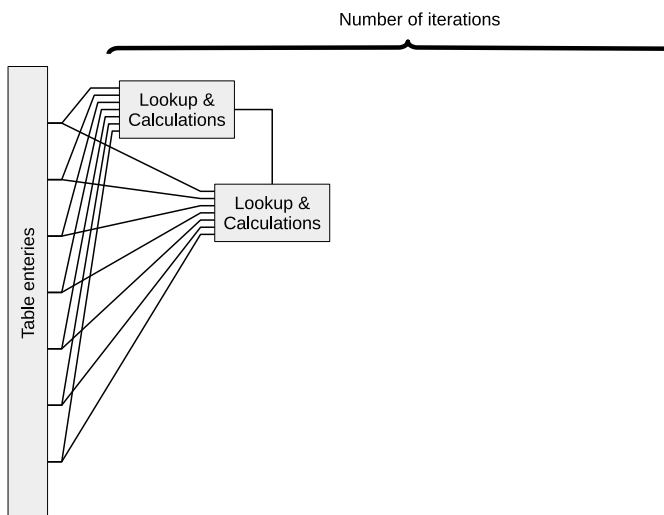
- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

2019-09-19

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME

Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!

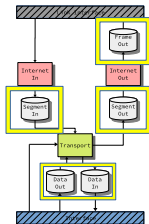
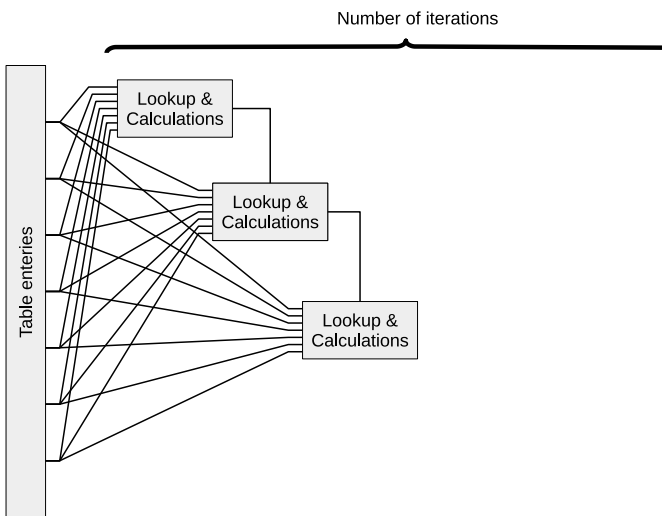
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME

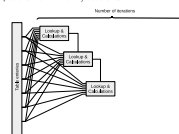
Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!

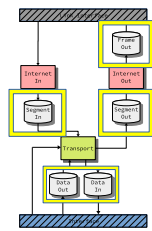
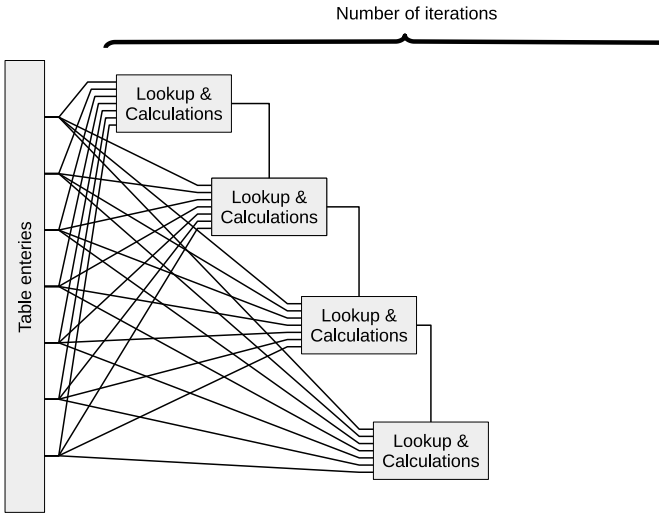
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME

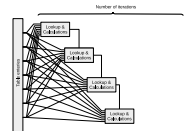
Implementation

Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!

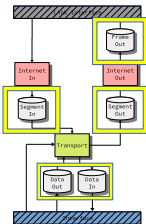
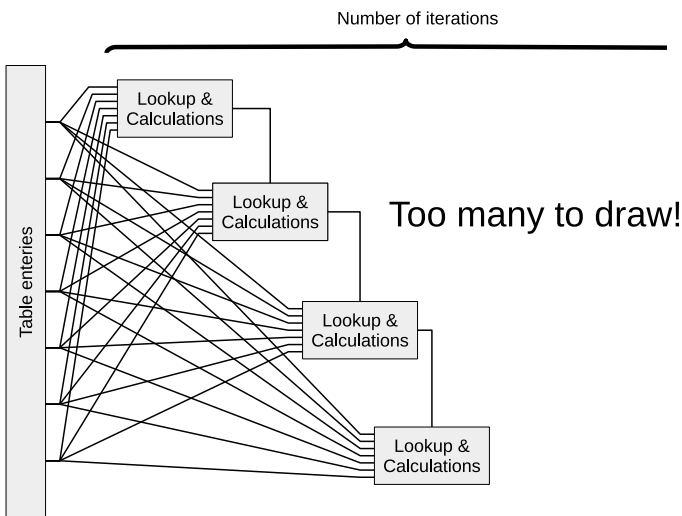
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Buffers

Some problems with the memory dictionaries!



TCP/IP in hardware using SME

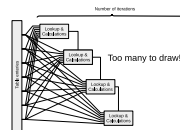
└ Implementation

└ Implementation

Implementation

Buffers

Some problems with the memory dictionaries!



Overflow!

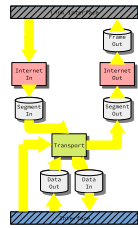
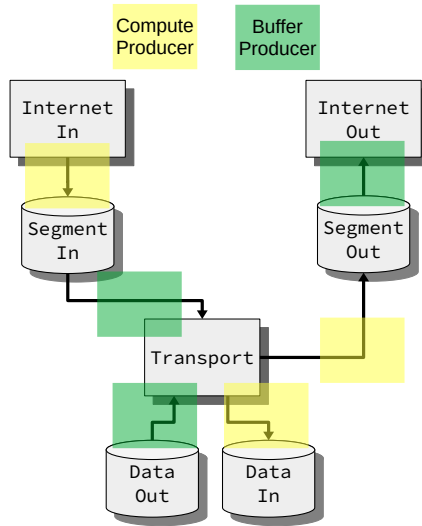
kan læses ved:

- Kør løkken en gang per clock
- Brug en anden model end en linked list, måske et fast offset?

Implementation

Interface signal protocol

Identifying the scenarios



TCP/IP in hardware using SME

Implementation

Implementation

Implementation
Interface signal protocol
Identifying the scenarios



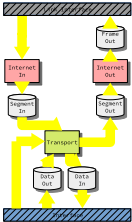
Data skal overføres hurtigst muligt, og det må ikke gå tabt

2 scenarier: fra "compute" til buffer, og omvendt

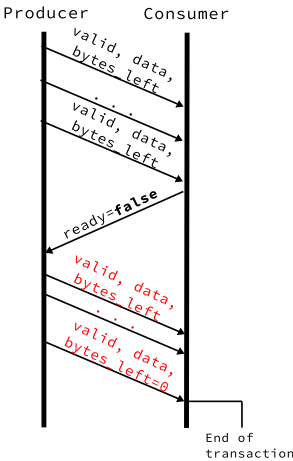
- CP kan ikke vente
- BP har stor buffer, og consumer starter transaktion

Implementation

Interface signal protocol



Compute-Producer (CP)



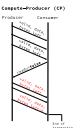
2019-09-19 TCP/IP in hardware using SME

└ Implementation

└ Implementation

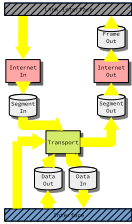
Implementation

Interface signal protocol

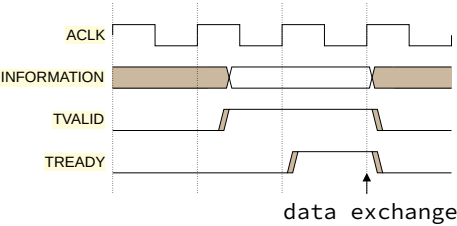


Implementation

Interface signal protocol



Buffer-Producer: Inspired by AXI4



TCP/IP in hardware using SME

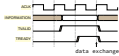
└ Implementation

└ Implementation

Implementation

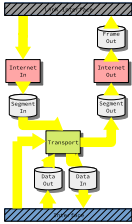
Interface signal protocol

Buffer-Producer: Inspired by AXI4

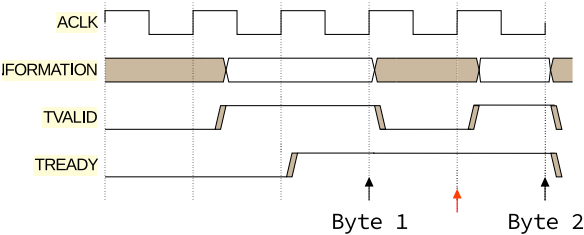


Implementation

Interface signal protocol



Streaming consecutive bytes can be a challenge!



TCP/IP in hardware using SME

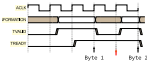
└ Implementation

└ Implementation

Implementation

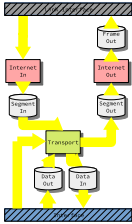
Interface signal protocol

Streaming consecutive bytes can be a challenge!



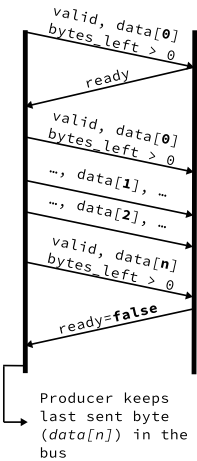
Implementation

Interface signal protocol



Buffer-Producer (BP)

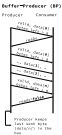
Producer Consumer



2019-09-19 TCP/IP in hardware using SME
└ Implementation

└ Implementation

Implementation
Interface signal protocol



1

Introduction

2

Implementation

3

Evaluation

• Setup

• Test

• Validation

4

Discussion

5

Conclusion

6

Future Work

7

Questions

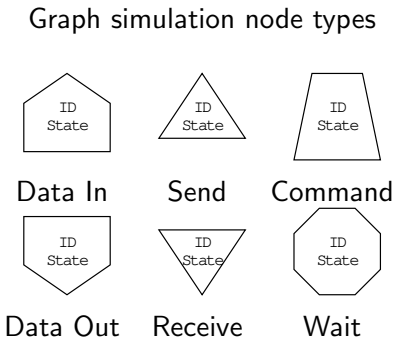
8

Demonstration

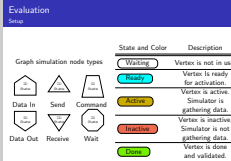
Graph file simulation

- Full input - output
- Does not take latency between packets into account
- Simplifies test cases

Definer send og receive bedre



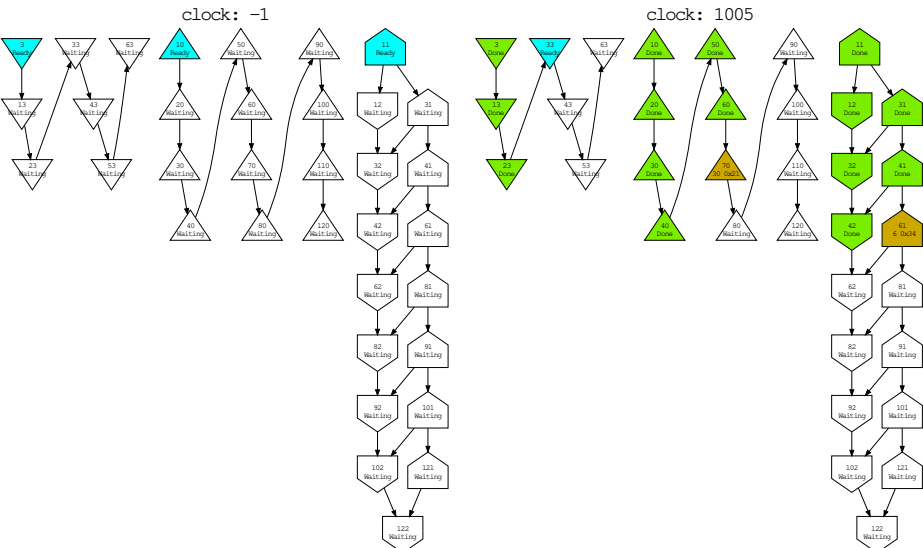
| State and Color | Description |
|-----------------|--|
| Waiting | Vertex is not in use. |
| Ready | Vertex is ready for activation. |
| Active | Vertex is active. Simulator is gathering data. |
| Inactive | Vertex is inactive. Simulator is not gathering data. |
| Done | Vertex is done and validated. |



Hop til illustrationen på næste slide nå du snakker om det!

Evaluation

Setup



Mark Jan Jacobi & Jan Meznik (KU)

TCP/IP in hardware using SME

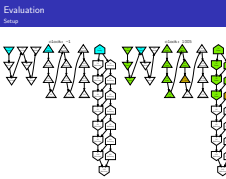
September 19, 2019

27 / 43

TCP/IP in hardware using SME

└ Evaluation

└ Evaluation



2019-09-19

Senario

- Real life scenario
- Test at high workloads
- Remove garbage
- Respond to packet
- Differ between concurrent connections

Fortæl kun om hvad vi vil have, ikke hvad vi har lavet af test

The test

- 17283 packets in total
- Two "sessions"
- 640*2 UDP packets that needs a response
- 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

TCP/IP in hardware using SME

└─Evaluation

└─Evaluation

Evaluation Test

The test

- 17283 packets in total
 - Two "sessions"
 - 640*2 UDP packets that needs a response
 - 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

Latency calculations:

n_D : The number of bytes in the data part of the protocol. This excludes both headers from transport and internet.

n_I : The internet header size.

n_T : The transport header size.

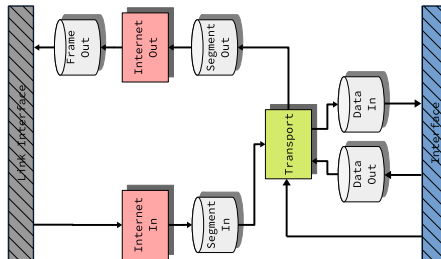
n : The total packet size.

From packet to user

$$6 + n_I + 2n_T + 3n_D$$

From user to packet

$$8 + 2n_I + 3n_T + 4n_D$$



TCP/IP in hardware using SME

└ Evaluation

└ Evaluation

Evaluation

Validation

Latency calculations:

n_D : The number of bytes in the data part of the protocol. This excludes both headers from transport and internet.

n_I : The internet header size.

n_T : The transport header size.

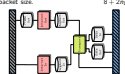
n : The total packet size.

From packet to user

$$6 + n_I + 2n_T + 3n_D$$

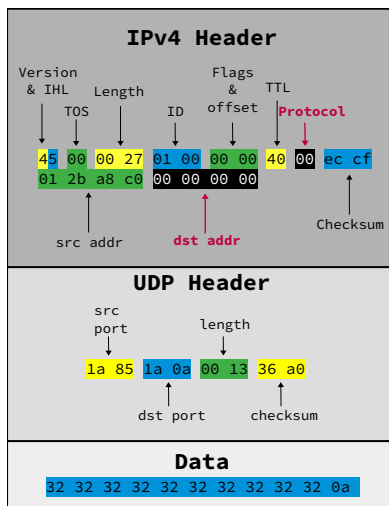
From user to packet

$$8 + 2n_I + 3n_T + 4n_D$$



Bufferen kan ikke videresende data direkte, da den skal gemme segmentet først

Outgoing packet validation:



TCP/IP in hardware using SME

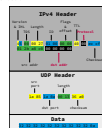
└ Evaluation

└ Evaluation

Evaluation

Validation

Outgoing packet validation:



Protocol ikke sat korrekt, destination ip ikke sat korrekt

Table of Contents

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion

6 Future Work

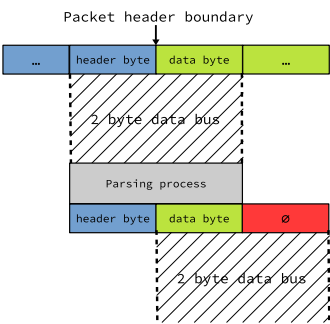
7 Questions

8 Demonstration

Estimated performance:

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

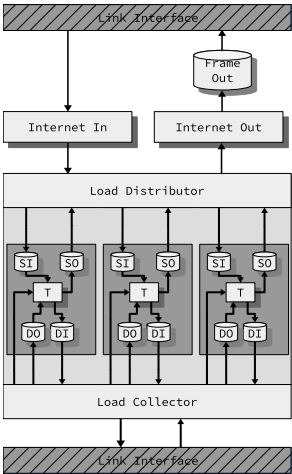
Improving the performance:



Estimated performance:

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

Improving the performance:



TCP/IP in hardware using SME

└ Discussion

└ Discussion

Discussion

Estimated performance

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

Improving the performance:

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion

6 Future Work

7 Questions

8 Demonstration

Conclusion

- Many design alterations, **layered** design worked best
- All 17283 ingoing packets parsed correctly
- Few easily fixable errors in outgoing packets
- SME was great for implementation, albeit with a few small errors and bugs

TCP/IP in hardware using SME

└ Conclusion

└ Conclusion

Conclusion

- Many design alterations, layered design worked best
- All 17283 ingoing packets parsed correctly
- Few easily fixable errors in outgoing packets
- SME was great for implementation, albeit with a few small errors and bugs

Design: Distribueret hukommelse er effektiv og giver funktionalitet

SME: C# nemt og hurtigt

Simulationen gjorde udviklingen hurtigere

Manglende features (structs), manglende library

1 Introduction

2 Implementation

3 Evaluation

4 Discussion

5 Conclusion

6 Future Work

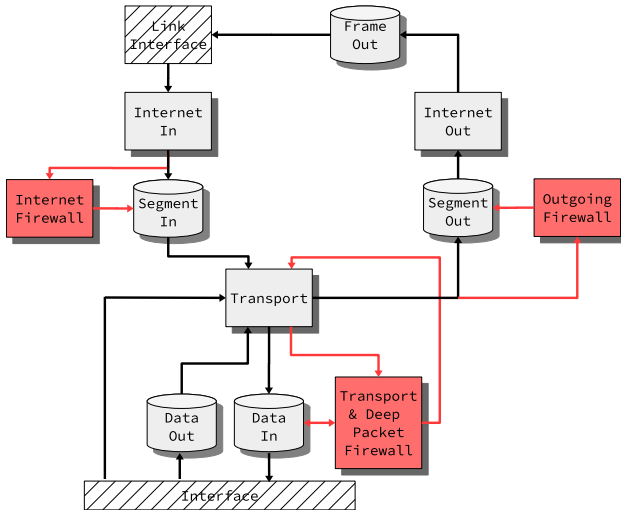
- Firewall
- TCP

7 Questions

8 Demonstration

Future Work

Firewall

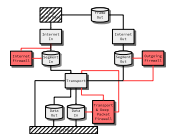


TCP/IP in hardware using SME

└ Future Work

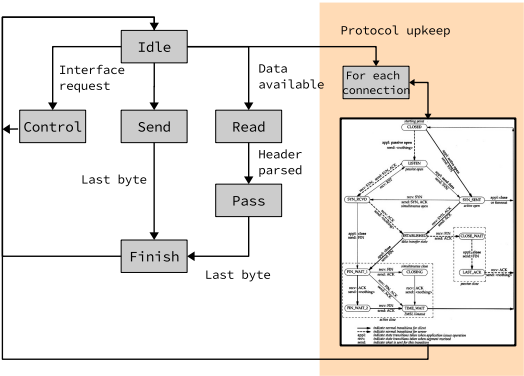
└ Future Work

Future Work
Firewall



Integration med buffere. Hvad ville det indebære

Implementing TCP



TCP/IP in hardware using SME

└ Future Work

└ Future Work

Future Work

TCP

Implementing TCP

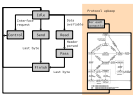


Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions
- 8 Demonstration

- [1] J. Weerasinghe, F. Abel, C. Hagleitner, and A. Herkersdorf. Disaggregated fpgas: Network performance comparison against bare-metal servers, virtual machines and linux containers. In *2016 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*, pages 9–17, Dec 2016. doi: 10.1109/CloudCom.2016.0018.

- 1

Introduction
- 2

Implementation
- 3

Evaluation
- 4

Discussion
- 5

Conclusion
- 6

Future Work
- 7

Questions
- 8

Demonstration

Demonstration

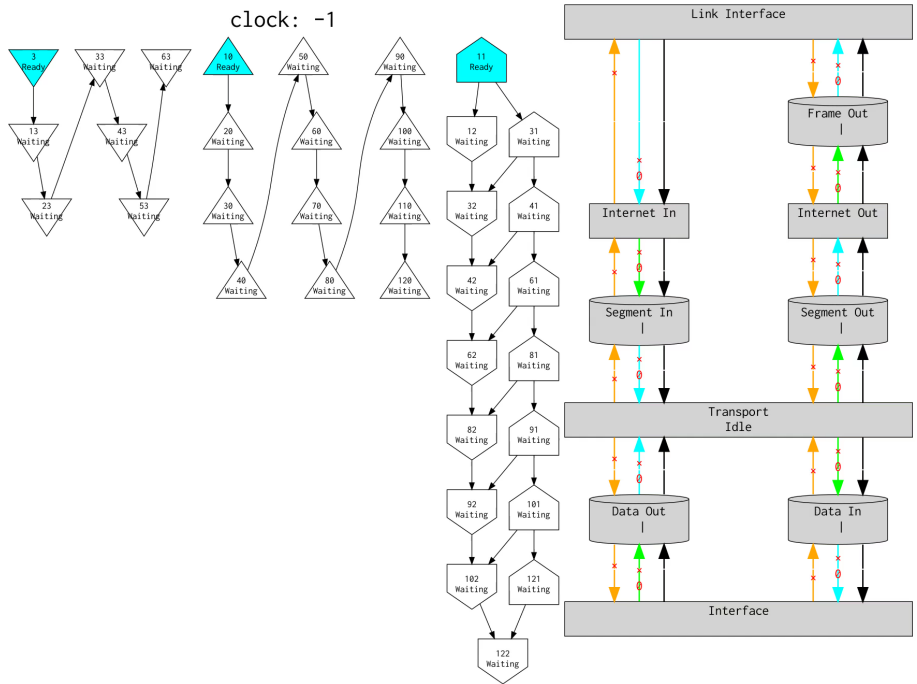
TCP/IP in hardware using SME

└─ Demonstration

└─ Demonstration

Demonstration

2019-09-19



TCP/IP in hardware using SME

Demonstration

