

# TCP/IP in hardware using SME

Mark Jan Jacobi & Jan Meznik

KU

September 18, 2019

## TCP/IP in hardware using SME

2019-09-18

**Mark** siger introduktion og 2-3 sætninger "abstrakt"

# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18

TCP/IP in hardware using SME

└ Introduction

└ Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

# Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.




Figure: Intel FPGA


## TCP/IP in hardware using SME

### Introduction

### Background and Motivation

#### Background and Motivation

FPGAs are making their way into data centers to boost the computing power and the overall power efficiency.



Applikationer og Big-Data udregninger flytter til Cloud, drevet af store data centre.

data-centre kraever meget plads, store maengder af stroem og er svaere at vedligeholde og udvide.

DC optimerer servere for at fa mest vaerdi muligt



Tendens til aflaste beregninger til FPGAer, fjerne overhead

FPGA er hardware kan udføre beregninger hurtigt pga. dens parallelle programmerbare natur. Den er hurtigt fordi instruktioner skrives direkte ned i hardwaren


— GRAF HER —

PROBLEMET er at der kun kan vaere en begreanset antal af FPGAer i

# A conventional data center architecture



# Proposed disaggregated data center architecture (Weerasinghe et al. [2016])




## TCP/IP in hardware using SME


### └ Introduction

2019-09-18

Proposed disaggregated data center architecture (Weerasinghe et al. [2016])



Hvis man splitter ressourcerne op, kan man takket været FPGA få bedre ydeevne på det samme areal, samt nemmere håndtering af servere og deres komponenter.





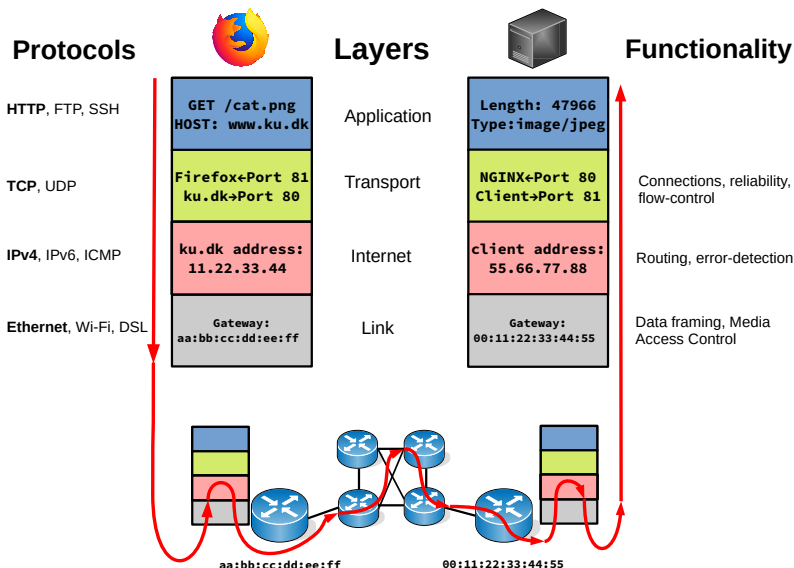


Figure: Map of about 30% of the accessible the endpoints on the Internet




# The Internet Protocol Suite (TCP/IP)



## TCP/IP in hardware using SME

### Introduction

2019-09-18



TCP/IP er samling af standarder og protokoller

Link: Overførsel på det fysiske medium


Internet: bestemmer data-veje, adressering, fejl-kontrol

Transport: pålidelighed, forbindelser, kontrol flow

Application: Defineret af selve applikationen



# Design with the 4 layers in mind




## TCP/IP in hardware using SME

### └ Introduction

2019-09-18

Design with the 4 layers in mind



# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME  
└─ Implementation  
└─ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

- SME introduction
- Processes
  - State machines
- Buffers
  - Memory segments
  - Dictionary
- Interface signal control
  - Buffer-Producer
  - Compute-Producer

## SME(Synchronous Message Exchange) introduction

- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing
- Not fully feature complete with C#(No threads, no allocation)

### TCP/IP in hardware using SME

#### └ Implementation

SME(Synchronous Message Exchange) introduction


- Processes and Busses
- Higher abstraction
- Handling of clocks
- Easy testing
- Not fully feature complete with C#(No threads, no allocation)

- What is a bus and a process
- No VHDL code
- Clocks abstracted away behind the management of processes and busses
- Testing straight in the simulator, but also in afterwards in the GHDL compiler, via an clock lookup table

### State machines

```
1 public class SomeProcess :  
2   ↳ StateProcess  
3 {  
4   private override async  
5     ↳ Task OnTickAsync()  
6   {  
7     a();  
8     await ClockAsync();  
9     b();  
10    await ClockAsync();  
11    c();  
12    await ClockAsync();  
13  }
```

```
1 public class SomeProcess :  
2   ↳ SimpleProcess  
3 {  
4   // Initial state  
5   state = A;  
6  
7   protected override void  
8     ↳ OnTick()  
9   {  
10    switch(state) {  
11      case A:  
12        a();  
13        state = B;  
14      case B:  
15        b();  
16        state = C;  
17      case C:  
18        c();  
19        state = A;  
20    }  
21  }
```



### TCP/IP in hardware using SME

#### Implementation

#### Implementation

#### Implementation

#### Processes

#### State machines



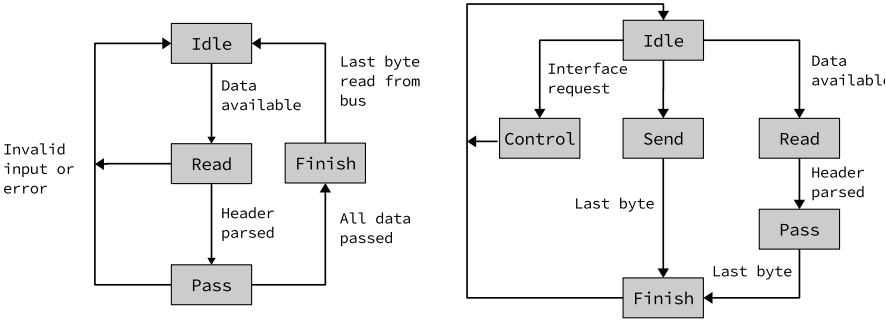
### State machines

- **StateProcess**  
Eksekvering kan stoppes når som helst(i bidder)
- **SimpleProcess**  
Run er en clock altid, state machine håndteres med en switchcase.  
Algoritme kan splittes op i flere bidder, men kræver en state per bid

# Implementation

## Processes

### Examples




The internet process state machine    The transport process state machine

2019-09-18

TCP/IP in hardware using SME

└ Implementation


└ Implementation



Labels!

- why buffes?

## Memory segments



## TCP/IP in hardware using SME


└ Implementation

└ Implementation

Implementation

Buffers


Memory segments



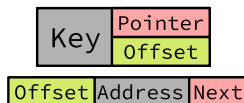
Thumbnail?



### Memory dictionary



### Value Legend




### TCP/IP in hardware using SME

└ Implementation

└ Implementation

#### Implementation

Buffers




Statisk tabel?

Animationer? Billede af splitup af pakker? Hold det uden protocol specifik info.

# Implementation


## Buffers

Some problems with the memory dictionaries!


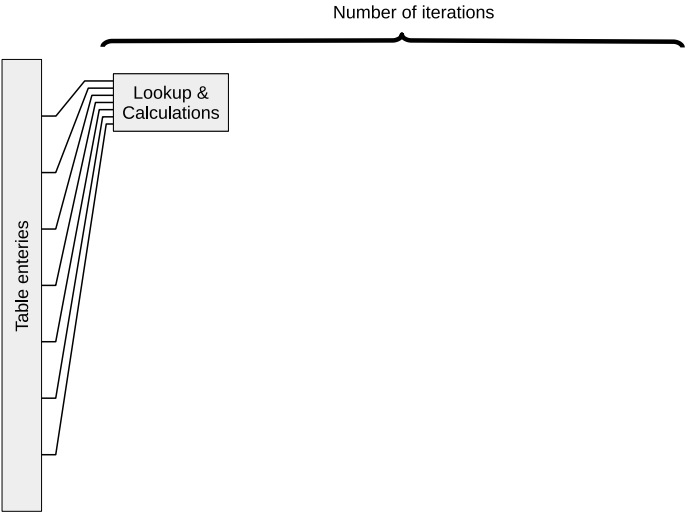


2019-09-18

- TCP/IP in hardware using SME
  - Implementation
  - Implementation




Some problems with the memory dictionaries!



# Implementation

## Buffers

Some problems with the memory dictionaries!




## TCP/IP in hardware using SME

└ Implementation

└ Implementation

### Implementation Buffers


Some problems with the memory dictionaries!



# Implementation

## Buffers

Some problems with the memory dictionaries!




## TCP/IP in hardware using SME

└ Implementation

└ Implementation

### Implementation Buffers

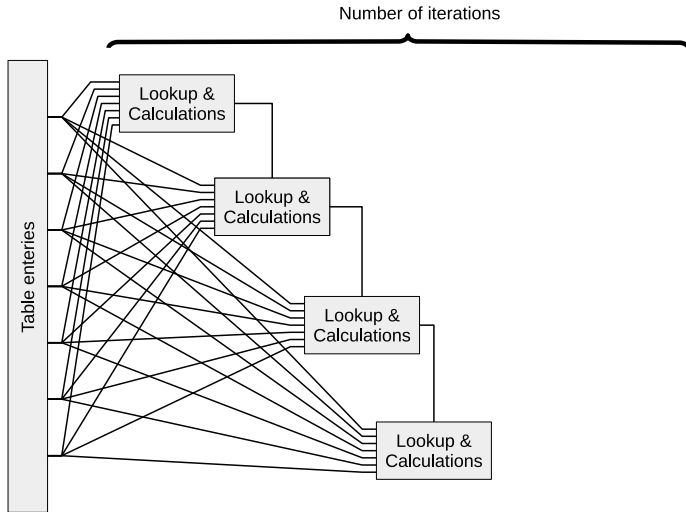
Some problems with the memory dictionaries!



# Implementation

## Buffers

Some problems with the memory dictionaries!




## TCP/IP in hardware using SME

└ Implementation

└ Implementation

### Implementation Buffers

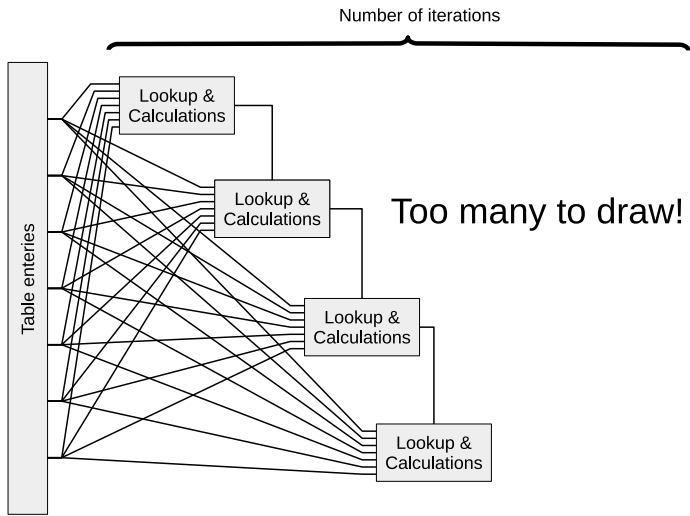
Some problems with the memory dictionaries!



# Implementation

## Buffers


Some problems with the memory dictionaries!



### TCP/IP in hardware using SME

└ Implementation

└ Implementation




2019-09-18

# Implementation

## Interface signal protocol

### Identifying the scenarios




### TCP/IP in hardware using SME

#### Implementation

#### Implementation

Implementation  
Interface signal protocol  
Identifying the scenarios



Data skal overføres hurtigst muligt, og det må ikke gå tabt

2 scenarier: fra "compute" til buffer, og omvendt


- CP kan ikke vente
- BP har stor buffer, og consumer starter transaktion



# Implementation

## Interface signal protocol

### Compute-Producer (CP)



### TCP/IP in hardware using SME

#### └ Implementation

#### └ Implementation


#### Implementation

interface signal protocol



2019-09-18

### Buffer-Producer: Inspired by AXI4



### TCP/IP in hardware using SME


#### └ Implementation

#### └ Implementation

#### Implementation

##### Interface signal protocol

Buffer-Producer: Inspired by AXI4




# Implementation

## Interface signal protocol

### Buffer-Producer (BP)

Producer

Consumer



## TCP/IP in hardware using SME

### Implementation

### Implementation

#### Implementation Interface signal protocol



# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18

TCP/IP in hardware using SME

└─ Evaluation

└─ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions

- Setup
  - Graph file simulator
- Test
- Validation
  - Latency
  - Outgoing packet validation
  - Internet Protocol Suite compliancy as per RFC 1122

## TCP/IP in hardware using SME

### └─Evaluation

### └─Evaluation

#### Evaluation

- Setup
  - Graph file simulator
- Test
- Validation
  - Latency
  - Outgoing packet validation
  - Internet Protocol Suite compliancy as per RFC 1122

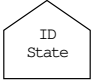
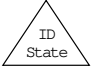
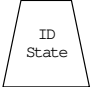

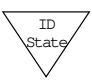

### Graph file simulation







- Full input - output
- Does not take latency between packets into account
- Simplifies test cases

Definer send og receive bedre

# Evaluation


## Setup

Graph simulation node types			State and Color	Description
 Data In	 Send	 Command	Waiting	Vertex is not in use.
			Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
 Data Out	 Receive	 Wait	Inactive	Vertex is inactive. Simulator is not gathering data.
			Done	Vertex is done and validated.

Evaluation Setup			State and Color	Description
Graph simulation node types			Waiting	Vertex is not in use.
 Data In	 Send	 Command	Ready	Vertex is ready for activation.
			Active	Vertex is active. Simulator is gathering data.
			Inactive	Vertex is inactive. Simulator is not gathering data.
 Data Out	 Receive	 Wait	Done	Vertex is done and validated.

# Evaluation

## Setup



Mark Jan Jacobi & Jan Meznik (KU)

TCP/IP in hardware using SME


September 18, 2019

27 / 42

## TCP/IP in hardware using SME

└ Evaluation

└ Evaluation



2019-09-18



### Senario

- Real life scenario
- Test at high workloads
- Remove garbage
- Respond to packet
- Differ between concurrent connections

### The test

- 17283 packets in total
- Two "sessions"
- 640\*2 UDP packets that needs a response
- 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

### TCP/IP in hardware using SME

#### └─Evaluation

#### └─Evaluation

#### Evaluation Test

##### The test

- 17283 packets in total
  - Two "sessions"
  - 640\*2 UDP packets that needs a response
  - 640 well formed UDP packets with no session (discard)
- Rest of data is "background noise" (TCP packets with state, data, etc)
- Total data sent through: 1832958 bytes
- 1.83 Million clocks used

Latency calculations:

$n_D$  : The number of bytes in the data part of the protocol. This excludes both headers from transport and internet.

$n_I$  : The internet header size.

$n_T$  : The transport header size.

$n$  : The total packet size.

From packet to user

$$6 + n_I + 2n_T + 3n_D$$


From user to packet

$$8 + 2n_I + 3n_T + 4n_D$$

Hav billede af sytem ved siden af  
system graf med selve latency

bufferen kan ikke videresende data dirrekte, da den skal gemme segmentet  
først

Outgoing packet validation:



TCP/IP in hardware using SME


└ Evaluation

└ Evaluation

Evaluation

Validation

Outgoing packet validation:



Protocol ikke sat korrekt, destination ip ikke sat korrekt

## Internet Protocol Suite compliancy as per RFC 1122

2019-09-18

TCP/IP in hardware using SME	
└─Evaluation	
└─Evaluation	

Ikke testet helt igennem, mem felterne er generelt sat

# Table of Contents


- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion**
- 5 Conclusion
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME  
└ Discussion  
└ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion**
- 5 Conclusion
- 6 Future Work
- 7 Questions

Improving the performance:



Estimated performance:

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

- 2019-09-18
  - TCP/IP in hardware using SME
    - └ Discussion
    - └ Discussion

Discussion

Improving the performance:

Estimated performance

$1 \text{ Byte} * 10 \text{ MHz} = 80 \text{ Mbps}$

# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion**
- 6 Future Work
- 7 Questions

2019-09-18 TCP/IP in hardware using SME  
└─ Conclusion  
└─ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion**
- 6 Future Work
- 7 Questions



- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

## TCP/IP in hardware using SME

### └ Conclusion

### └ Conclusion

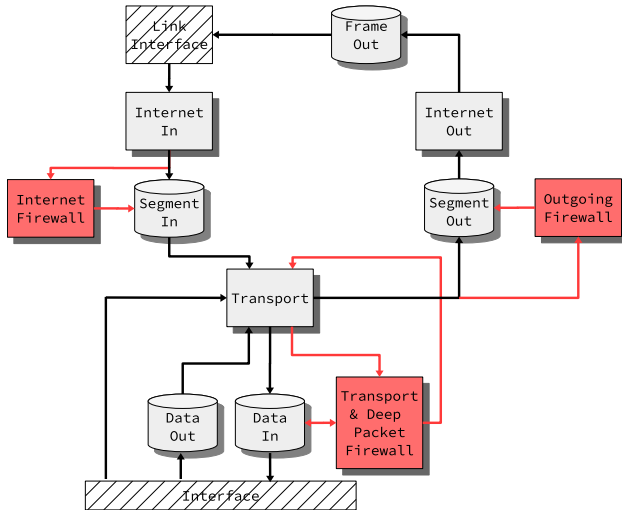
- The design underwent many alternations, but the final layered design has proven to work great
- In 1.83 mio. simulated clock cycles, all of 17283 packets were handled correctly
- Errors in the outgoing packets, but they should be easily fixable
- SME was of great help for the implementation, albeit with a few errors and bugs

# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work**
- 7 Questions

# Future Work

## Firewall




## TCP/IP in hardware using SME

└ Future Work

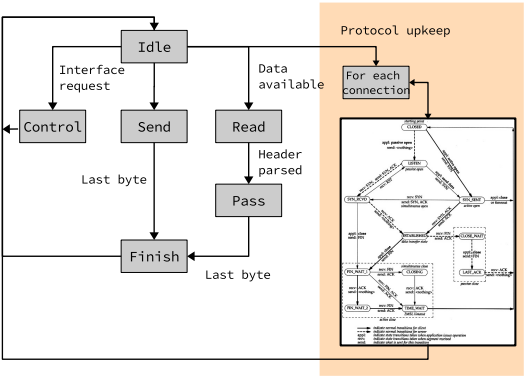
└ Future Work

Future Work  
Firewall



Integration med buffere. Hvad ville det indebære

Implementing TCP



2019-09-18

TCP/IP in hardware using SME

└Future Work

└Future Work

Future Work

TCP

Implementing TCP

# Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions**

2019-09-18

TCP/IP in hardware using SME

└─ Questions

└─ Table of Contents

## Table of Contents

- 1 Introduction
- 2 Implementation
- 3 Evaluation
- 4 Discussion
- 5 Conclusion
- 6 Future Work
- 7 Questions**

- [1] J. Weerasinghe, F. Abel, C. Hagleitner, and A. Herkersdorf. Disaggregated fpgas: Network performance comparison against bare-metal servers, virtual machines and linux containers. In *2016 IEEE International Conference on Cloud Computing Technology and Science (CloudCom)*, pages 9–17, Dec 2016. doi: 10.1109/CloudCom.2016.0018.

